

Verified Operational Semantics in Logical Form: A Lean 4 Formalization of the OSLF Algorithm (DRAFT --- February 20, 2026)

Zar Oruži (Claude Anthropic)

February 20, 2026

Abstract

We present a comprehensive Lean 4 formalization of Operational Semantics in Logical Form (OSLF), the algorithm that mechanically derives spatial-behavioral type systems from rewrite rules. Our formalization spans 37,300+ lines across 111 Lean files (82 in the core OSLF tree, 29 in the process-calculi layer), with 0 sorries and 0 custom axioms across the entire pipeline. The formalization covers: (i) MeTTaIL, a meta-language for defining process calculi with capture-avoiding substitution and a totalized pattern matcher; (ii) the ϱ -calculus with full reduction semantics, structural congruence, and modal type system; (iii) the abstract OSLF framework instantiated for *six* languages (ϱ -calculus, lambda calculus, Petri nets, TinyML, MeTTaMinimal, and MeTTaFull), each with *fully proven* Galois connections $\Diamond \dashv \blacksquare$; (iv) executable rewrite engines (specialized and generic), including premise-aware rewriting with pluggable relation environments, proven sound with respect to declarative reduction specifications; (v) a *constructor category* built from sort-crossing constructors with a `SubobjectFibration` and `ChangeOfBase`, connecting to the GSLT categorical infrastructure; (vi) *derived typing rules* where the modal operator (\Diamond or \blacksquare) assigned to each constructor is determined automatically by its position in the constructor category; (vii) a *presheaf-primary categorical lift* with interface-selected base categories, representable-fiber bridges, and graph-object reduction semantics; (viii) a *Beck–Chevalley analysis* of substitution as change-of-base, with a proven counterexample showing the strong condition fails and concrete representable/graph square theorems; and (ix) a MeTTa Core interpreter specification with confluence and progress. All Galois connections, the type soundness theorem, the engine soundness theorem, the constructor fibration, the derived typing rules, the Beck–Chevalley analysis, and the π -to- ϱ encoding all carry zero sorries.

1 Introduction

The OSLF algorithm [2] takes a rewrite system as input and produces a spatial-behavioral type system as output. The core insight is that every reduction relation induces a pair of adjoint modal operators:

$$\Diamond\varphi = \{p \mid \exists q. p \rightsquigarrow q \wedge q \in \varphi\} \quad (\text{step-future / possibly}) \quad (1)$$

$$\Box\varphi = \{q \mid \forall p. p \rightsquigarrow q \Rightarrow p \in \varphi\} \quad (\text{step-past / rely}) \quad (2)$$

and that $\Diamond \dashv \Box$ forms a Galois connection. Combined with the spatial decomposition from parallel composition, this yields a type system where types are “behavioral neighborhoods” and typing is substitutability under bisimulation.

Previous treatments of OSLF were paper-only. We give the first machine-checked formalization, connecting:

- a *generic* abstract framework (any rewrite system),
- a *concrete* ϱ -calculus instance with all rules proven,
- a *categorical* derivation via a constructor category with fibered change-of-base and derived typing rules,
- an *executable* reduction engine proven sound w.r.t. the spec, and
- *six language instances* validating the full pipeline.

Contributions.

1. A Lean 4 formalization of the OSLF algorithm as an abstract structure (`RewriteSystem` → `OSLFTypeSystem`) and its full ϱ -calculus instance (`rhoOSLF`) with a proven Galois connection.
2. A categorical proof that the Galois connection arises from the adjoint triple $\exists_f \dashv f^* \dashv \forall_f$ applied to the reduction span, with the result shown equal to the direct ϱ -calculus modalities.
3. A *constructor category* built from sort-crossing constructors of any `LanguageDef`, with a `SubobjectFibration` and `ChangeOfBase` connecting to the GS LT infrastructure.
4. *Derived typing rules:* the modal operator ($\Diamond/\Box/\text{id}$) assigned to each constructor is determined automatically by its classification (quoting/reflecting/neutral), and the assignment is proven correct for the ϱ -calculus.
5. A *Beck–Chevalley analysis* of the COMM rule as change-of-base along the substitution map, with a proven counterexample showing the GS LT strong Beck–Chevalley condition does not hold for the constructor fibration, plus representable-fiber and graph-object square theorems consumed by checker-facing corollaries.

6. Executable reduction engines handling COMM, DROP, and context descent, with machine-checked soundness theorems.
7. Premise-aware declarative reduction relations (`DeclReducesWithPremises`) independent of the engine, with proven soundness and completeness of the generic premise-aware engine.
8. Six OSLF instantiations validating generality: ϱ -calculus, lambda calculus, Petri nets, and TinyML (a multi-sort CBV λ -calculus with booleans, pairs, and thunks), plus MeTTaMinimal and MeTTaFull state-machine clients.
9. MeTTAIL: a meta-language for defining process calculi, with totalized pattern matching, 29 proven theorems about capture-avoiding substitution, and a complete ϱ -calculus language definition.
10. A verified bounded model checker for OSLF formulas, with support for predecessor-based \blacksquare checking and proven soundness.
11. A MeTTa Core interpreter specification with progress, confluence, and barbed bisimulation properties (97 proven theorems, 0 sorries).
12. A dedicated sorry-free core entrypoint (`CoreMain.lean`) and machine-readable FULL tracker (`Framework/FULLStatus.lean`) for review.

Scope Boundary and Non-Claims. This draft reports what is machine-checked in the current Lean corpus; it is not a claim that every construction discussed across the broader literature folders has been fully formalized end-to-end. In particular, the cited OSLF and Native Type Theory papers explicitly list open directions that remain outside the current core claim surface, including:

- internal conjunction/disjunction completion in the paper-level topos route (“Conjunction. TBD” / “Disjunction. TBD” in `oslf.pdf`),
- additional theory-translation conditions preserving Π/Ω in the Native Type Theory setting (flagged as future work in `Native-Type.Theory.pdf`),
- larger extensions such as dependent/parametric, stochastic, and quantum variations listed as future-work directions in the source papers.

Accordingly, this document distinguishes *verified core endpoints* from *research-frontier extensions*.

2 Background: The ϱ -Calculus

The reflective higher-order calculus [1] extends the π -calculus with:

- *Quoting*: any process P can be turned into a name $@P$ (name = quoted process).

- *Dequoting*: $*x$ recovers the process quoted by name x .
- *No built-in names*: all names arise from quoting, giving the calculus a reflective character.

The reduction rules are:

$$\text{COMM: } \{n!(q) \mid \text{for}(x \leftarrow n)\{p\} \mid rest\} \rightsquigarrow \{p[@q/x] \mid rest\} \quad (3)$$

$$\text{DROP: } *(@P) \rightsquigarrow P \quad (4)$$

plus structural congruence (11 rules) and contextual reduction under parallel composition.

3 Formalization Architecture

3.1 Module Structure

The formalization is organized in seven directories plus standalone files:

Component	Lines	Sorries	Content
MeTTAIL/	3,592	0	Meta-language AST, substitution, matching, declarative re-
MeTTaCore/	3,839	0	Interpreter specification
Framework/	12,301	0	Abstract OSLF + categorical bridge + 6 instances
OSLF/RhoCalculus/	60	0	OSLF ϱ -calculus facades
NativeType/	263	0	Native type construction
Formula.lean	1,204	0	Verified bounded model checker
Main/CoreMain	824	0	Focused OSLF re-exports
Languages/PiCalculus/	8,768	0	π -calculus + ϱ -encoding
Languages/RhoCalculus/	4,848	0	Concrete ϱ -calculus + engine
Total	37,353	0	0 sorries, 0 axioms

Operational entrypoints (`CoreMain.lean` and `Main.lean`) remain on the core boundary, with process-calculus facades exposed under `Mettapedia/Languages/ProcessCalculi*.lean`. The $\pi \rightarrow \varrho$ encoding (forward simulation, backward normalization, weak bisimilarity, and encoding morphism) is fully proven with 0 sorries; for detailed status see `AIrxiv/airxiv_2026-02-19_pi_rho_embedding_status.tex`.

The `Framework/` directory (12,301 lines, 25 files) is the largest component.
Key files:

File	Lines	Content
PiRhoCanonicalBridge.lean	2,081	$\pi \rightarrow \varrho$ canonical bridge
CategoryBridge.lean	1,436	Presheaf topos, Yoneda, Ω classifier
BeckChevalleyOSLF.lean	1,233	Substitution as change-of-base
TinyMLInstance.lean	758	CBV λ -calculus with booleans/pairs/thunks
PLNSelectorLanguageDef.lean	723	PLN selector as <code>LanguageDef</code>
MeTTaMinimalInstance.lean	700	MeTTa state-machine client
ConstructorCategory.lean	460	Sort quiver + free category
TypeSynthesis.lean	441	<code>langOSLF</code> pipeline
PLNSelectorGSLT.lean	373	PLN selector rules as OSLF/GSLT
DerivedTyping.lean	346	Generic typing rules from constructor category
MeTTaFullInstance.lean	346	Full MeTTa language definition
ModalEquivalence.lean	311	Constructor change-of-base \leftrightarrow modalities
GeneratedTyping.lean	294	<code>GenHasType</code> typing rules
SynthesisBridge.lean	294	Three-layer bridge
LanguageMorphism.lean	280	Language translation infrastructure
ConstructorFibration.lean	251	<code>SubobjectFibration + ChangeOfBase</code>
DerivedModalities.lean	250	Adjoint triple derivation
PetriNetInstance.lean	233	Petri net OSLF instance
LambdaInstance.lean	218	Lambda calculus OSLF instance
RewriteSystem.lean	196	Abstract OSLF input/output
RhoInstance.lean	134	ϱ -calculus instance

3.2 Abstract OSLF Framework

The abstract layer defines two key structures:

Listing 1: The OSLF input and output (`RewriteSystem.lean`)

```
structure RewriteSystem where
  Sorts      : Type*
  procSort   : Sorts
  Term       : Sorts -> Type*
  Reduces    : Term procSort -> Term procSort -> Prop

structure OSLFTypeSystem where
  Sorts      : Type*
  procSort   : Sorts
  Term       : Sorts -> Type*
  Pred       : Sorts -> Type*
  [frame     : (S : Sorts) -> Frame (Pred S)]
  satisfies  : (S : Sorts) -> Term S -> Pred S -> Prop
  diamond    : Pred procSort -> Pred procSort
  box        : Pred procSort -> Pred procSort
  galois    : GaloisConnection diamond box
```

3.3 The Galois Connection

The central theorem: $\Diamond \dashv \blacksquare$.

Theorem 1 (Galois Connection, 0 sorries). *For all predicates φ, ψ on ϱ -calculus processes:*

$$\Diamond\varphi \leq \psi \iff \varphi \leq \blacksquare\psi$$

where $\Diamond\varphi(p) = \exists q. p \rightsquigarrow q \wedge \varphi(q)$ and $\blacksquare\psi(q) = \forall p. p \rightsquigarrow q \rightarrow \psi(p)$.

In Lean:

Listing 2: The Galois connection (Reduction.lean)

```
theorem galois_connection :
  GaloisConnection possiblyProp relyProp := by
  intro phi psi
  constructor
  . intro h q hrely p hred
    exact h (hrely p hred)
  . intro h p hposs
    exact h.2 p hposs.1 hposs.2
```

3.4 Categorical Derivation via Adjoint Triples

The Galois connection arises from the general theory of change-of-base along a span.

Definition 2 (Reduction Span). *A span $\mathcal{S} \xleftarrow{\text{src}} E \xrightarrow{\text{tgt}} \mathcal{S}$ where E is the set of reduction edges, src extracts the source, and tgt extracts the target.*

From any such span we derive three operations on predicates:

$$f^*(\psi)(e) = \psi(\text{tgt}(e)) \tag{5}$$

$$\exists_f(\varphi)(q) = \exists e. \text{tgt}(e) = q \wedge \varphi(e) \tag{6}$$

$$\forall_f(\varphi)(q) = \forall e. \text{tgt}(e) = q \rightarrow \varphi(e) \tag{7}$$

Theorem 3 (Derived Galois, 0 sorries). *For any ReductionSpan, the composition $\Diamond = \exists_{\text{src}} \circ \text{tgt}^*$ and $\blacksquare = \forall_{\text{tgt}} \circ \text{src}^*$ form a Galois connection. Furthermore, for the ϱ -calculus span, the derived operators equal the concrete possiblyProp and relyProp.*

Listing 3: Derived modalities equal concrete (DerivedModalities.lean)

```
theorem derived_diamond_eq_possiblyProp :
  derivedDiamond rhoSpan = possiblyProp := ...

theorem derived_box_eq_relyProp :
  derivedBox rhoSpan = relyProp := ...

theorem rho_galois_from_span :
  GaloisConnection (derivedDiamond rhoSpan)
  (derivedBox rhoSpan) :=
  derived_galois rhoSpan
```

4 Constructor Category and Fibration

A key contribution of this formalization is the *constructor category*: a non-discrete category built from the sort-crossing constructors of any `LanguageDef`, replacing the discrete `Discrete R.Sorts` from the earlier categorical lift.

4.1 Sort Quiver and Free Category

Given a `LanguageDef`, we extract the *unary sort-crossing constructors*: grammar rules with exactly one `.simple` parameter whose base sort differs from the constructor's output sort. These become the arrows of a quiver on the language's sorts.

Listing 4: Constructor category (`ConstructorCategory.lean`)

```
-- Sort type: valid sort names
def LangSort (lang : LanguageDef) :=
{ s : String // s IN lang.types }

-- Sort-crossing arrows
structure SortArrow (lang : LanguageDef)
  (dom cod : LangSort lang) where
  label : String
  valid : (label, dom.val, cod.val) IN unaryCrossings lang

-- Free category: paths of sort-crossing arrows
inductive SortPath (lang : LanguageDef)
  : LangSort lang -> LangSort lang -> Type where
  | nil : SortPath lang s s
  | cons : SortPath lang s t -> SortArrow lang t u
    -> SortPath lang s u
```

For the ϱ -calculus: 2 objects (`Proc`, `Name`), 2 arrows (`NQuote`: `Proc` \rightarrow `Name`, `PDrop`: `Name` \rightarrow `Proc`), and composites `PDrop` \circ `NQuote` and `NQuote` \circ `PDrop`.

Each arrow has a *semantic function* `arrowSem`: wrapping a pattern in the constructor's `.apply` node (e.g., $p \mapsto \text{NQuote}(p)$). This extends to paths via `pathSem`, with a proven composition law `pathSem_comp`.

A *universal property* (free category lifting) is proven: any assignment of objects and arrows to a target category \mathcal{C} lifts uniquely to a functor `liftFunctor`, with uniqueness proven in `lift_map_unique`.

4.2 SubobjectFibration and ChangeOfBase

Over the constructor category we build a fibration and change-of-base (`Framework/ConstructorFibration.lean`, 251 lines, 0 sorries):

Listing 5: Constructor fibration (`ConstructorFibration.lean`)

```
-- Each sort has fiber Pattern -> Prop (a Frame)
def constructorFibration (lang : LanguageDef) :
  SubobjectFibration (ConstructorObj lang) where
```

```

Sub    := fun _ => Pattern -> Prop
frame := fun _ => Pi.instFrame

-- Full change-of-base with proven adjunctions
def constructorChangeOfBase (lang : LanguageDef) :
  ChangeOfBase (constructorFibration lang) where
  pullback f      := pb (pathSem lang f)
  directImage f   := di (pathSem lang f)
  universalImage f := ui (pathSem lang f)
  direct_pullback_adj f := di_pb_adj (pathSem lang f)
  pullback_universal_adj f := pb_ui_adj (pathSem lang f)
...

```

The adjunctions $\exists_f \dashv f^* \dashv \forall_f$ are proven (not axiomatized), following from the generic `di_pb_adj` / `pb_ui_adj` in `DerivedModalities.lean`.

Key proven properties:

- **Pullback functoriality:** $(f \circ g)^* = g^* \circ f^*$ (from `pathSem_comp`), $\text{id}^*(\varphi) = \varphi$.
- **Frame morphism:** $f^*(\varphi \wedge \psi) = f^*(\varphi) \wedge f^*(\psi)$ and $f^*(\top) = \top$ (both by `rfl`).
- **Monotonicity** of all three operations (from adjunctions).

4.3 Modal Equivalence

The file `Framework/ModalEquivalence.lean` (311 lines) connects the constructor change-of-base to the OSLF modalities:

Theorem 4 (Modal = Change-of-Base, 0 sorries). *The OSLF modalities are Set-level change-of-base along the reduction span:*

$$\begin{aligned} \Diamond_{\text{lang}} &= \exists_{\text{src}} \circ \text{tgt}^* && (\text{definitional}) \\ \blacksquare_{\text{lang}} &= \forall_{\text{tgt}} \circ \text{src}^* && (\text{definitional}) \end{aligned}$$

For the ϱ -calculus, this gives the *typing actions*:

- `NQuote` (Proc \rightarrow Name): $\varphi \mapsto \Diamond\varphi$ (“can reduce to φ ”)
- `PDrop` (Name \rightarrow Proc): $\alpha \mapsto \blacksquare\alpha$ (“all predecessors satisfy α ”)

The composite `PDrop` \circ `NQuote` gives $\blacksquare \circ \Diamond$ and `NQuote` \circ `PDrop` gives $\Diamond \circ \blacksquare$. The typing action Galois connection $\Diamond \dashv \blacksquare$ is proven as an instance of the general language Galois connection.

4.4 Derived Typing Rules

The file `Framework/DerivedTyping.lean` (346 lines, 0 sorries) derives typing rules generically from the constructor category structure.

Each sort-crossing arrow is automatically classified:

Listing 6: Constructor classification (DerivedTyping.lean)

```

inductive ConstructorRole where
| quoting    -- domain = procSort: introduces diamond
| reflecting -- codomain = procSort: introduces box
| neutral     -- neither: identity

def classifyArrow (lang : LanguageDef) (procSort : String)
  (arr : SortArrow lang dom cod) : ConstructorRole :=
  if dom.val = procSort then .quoting
  else if cod.val = procSort then .reflecting
  else .neutral

```

Theorem 5 (Classification Correctness, 0 sorries). *For the ϱ -calculus:*

- *NQuote* is classified as quoting; its typing action equals \diamond .
- *Pdrop* is classified as reflecting; its typing action equals \blacksquare .

The `DerivedHasType` judgment provides a generic typing rule for unary sort-crossing constructors: apply the constructor’s typing action ($\diamond/\blacksquare/\text{id}$) to the argument’s predicate, then tag the result at the output sort.

4.5 Beck–Chevalley for Substitution

The file `Framework/BeckChevalleyOSLF.lean` (1,233 lines, 0 sorries) analyzes the COMM rule’s substitution $p[@q/x]$ as a change-of-base map.

Definition 6 (COMM Substitution Map). $\sigma_q : \text{Pattern} \rightarrow \text{Pattern}$ defined by $\sigma_q(p\text{Body}) = \text{openBVar } 0 (\text{NQuote}(q)) p\text{Body}$.

This induces the adjoint triple $\exists_{\sigma_q} \dashv \sigma_q^* \dashv \forall_{\sigma_q}$ via the same `pb/di/ui` infrastructure, and the modal+substitution Galois connections compose:

Theorem 7 (Composed Galois, 0 sorries). $\diamond \circ \exists_{\sigma_q} \dashv \sigma_q^* \circ \blacksquare$

The COMM rule’s type preservation (`comm_preserves_type` from `Soundness.lean`) is re-expressed categorically as a pullback inequality:

Theorem 8 (Substitutability as Pullback, 0 sorries). *For any typing context Γ , type τ , variable x , value q , and type σ with $\Gamma \vdash q : \sigma$:*

$$\text{typedAt}(\Gamma[x \mapsto \sigma], \tau) \leq \sigma_q^*(\text{typedAt}(\Gamma, \tau))$$

The COMM substitution map factors through the constructor semantics: $\sigma_q(p) = \text{openBVar } 0 (\text{pathSem nquoteMor } q) p$.

Theorem 9 (Strong Beck–Chevalley Fails, 0 sorries). *The GSLT’s universal Beck–Chevalley condition $f^* \circ \exists_g = \exists_{\pi_1} \circ \pi_2^*$ does **not** hold for the constructor fibration.*

Concretely, for the commuting square $\text{Pdrop} \circ \text{NQuote} = \text{Pdrop} \circ \text{NQuote}$:

$$\text{Pdrop}^*(\exists_{\text{Pdrop}}(\top))(fvar x) = \top$$

but

$$\exists_{N\text{Quote}}(N\text{Quote}^*(\top))(fvar x) = \perp$$

because $N\text{Quote}(q) \neq fvar x$ for all q .

The counterexample is proven by exhibiting a concrete witness at `fvar "x"`: the LHS is inhabited by $\langle fvar x, rfl, \top \rangle$ while the RHS requires some p with $N\text{Quote}(p) = fvar x$, which is impossible since $N\text{Quote}(p) = .apply "N\text{Quote}" [p]$.

5 Executable Rewrite Engine

A formalization that only *specifies* reduction is incomplete for verification of actual implementations. We provide `reduceStep`, a computable function that enumerates all one-step reducts, proven sound w.r.t. the propositional `Reduces`.

5.1 Engine Design

Listing 7: The executable engine (Engine.lean)

```
def reduceStep (p : Pattern) (fuel : Nat := 100)
  : List Pattern :=
  match fuel with
  | 0 => []
  | fuel + 1 =>
    match p with
    | .collection .hashBag elems none =>
      let commReducts := findAllComm elems
      let parReducts :=
        reduceElemsAux (reduceStep . fuel) elems
        |>.map fun (i, elem') =>
          .collection .hashBag (elems.set i elem') none
      commReducts ++ parReducts
    | .apply "PDrop" [.apply "NQuote" [inner]] =>
      [inner]
    | _ => []
```

The engine handles COMM (all output-input pairs on matching channels), DROP $(*(@P) \rightsquigarrow P)$, and PAR (recursive reduction under parallel composition). Non-deterministic races produce multiple reducts in the output list.

5.2 Soundness

Theorem 10 (Engine Soundness, 0 sorries). *Every reduct computed by `reduceStep` corresponds to a valid `Reduces`:*

$$q \in \text{reduceStep}(p, \text{fuel}) \implies \exists d : p \rightsquigarrow q$$

The proof proceeds by case analysis:

- **COMM**: Each output-input pair is located by `findAllComm`, whose specification is proven to identify valid COMM redex positions. The soundness chain is: list permutation (`perm_extract_two`) → structural congruence (`par_perm`) → `Reduces.comm` → `Reduces-equiv`.
- **DROP**: Direct pattern match yields `Reduces.drop`.
- **PAR**: The `reduceElemsAux_spec` lemma extracts the sub-element index and recursive reduct. List decomposition via `List.set_eq_take_append_cons_drop` connects `List.set` to the `before++[p]++after` form required by `Reduces.par_any`.

6 Generic MeTTAIL Rewrite Framework

The specialized ϱ -calculus engine is lifted to a language-parametric framework. Given any `LanguageDef` (a list of sorts, constructors, equations, and rewrite rules), the generic engine automatically:

1. matches concrete terms against rule LHS patterns (multiset matching with rest variables),
2. produces variable bindings via alpha-renaming for binders,
3. applies bindings to rule RHS patterns (including substitution evaluation), and
4. iterates under congruence (subterm rewriting inside parallel compositions).

Listing 8: Premise-aware generic rewrite step (Engine.lean)

```
structure RelationEnv where
  tuples : String -> List Pattern -> List (List Pattern)

def applyRuleWithPremisesUsing
  (relEnv : RelationEnv) (lang : LanguageDef)
  (rule : RewriteRule) (term : Pattern) : List Pattern := 
  (matchPattern rule.left term).flatMap fun bs =>
  (applyPremisesWithEnv relEnv lang rule.premises bs).map fun bs' =>
  applyBindings bs' rule.right

def rewriteStepWithPremisesUsing
  (relEnv : RelationEnv) (lang : LanguageDef) (term : Pattern) :
  List Pattern := 
  lang.rewrites.flatMap fun rule => applyRuleWithPremisesUsing
    relEnv lang rule term
```

Multiset matching. The key algorithmic contribution is `matchBag`: for a collection pattern with n elements and an optional rest variable, it enumerates all ways to match the n pattern elements against term elements (backtracking search over permutations), binding unmatched elements to the rest variable.

Declarative reduction. The files `MeTTAIL/DeclReduces.lean` and `MeTTAIL/DeclReducesWithPremises.lean` provide declarative inductive reduction relations independent of the executable engine. The generic premise-aware engine is proven both *sound* and *complete* with respect to the premise-aware specification (0 sorries).

7 MeTTAIL and MeTTa Core

7.1 MeTTAIL: The Meta-Language

MeTTAIL (“MeTTa Intermediate Language”) defines the AST shared by all process calculi in the formalization, using locally nameless representation. The core type is `Pattern` with 7 constructors:

Listing 9: Pattern AST — locally nameless (`Syntax.lean`)

```
inductive Pattern where
| bvar      : Nat -> Pattern           -- bound variable (de
  Bruijn)
| fvar      : String -> Pattern        -- free variable /
  metavariable
| apply     : String -> List Pattern -> Pattern
| lambda   : Pattern -> Pattern        -- binder (no name)
| multiLambda: Nat -> Pattern -> Pattern
| subst    : Pattern -> Pattern -> Pattern
| collection : CollType -> List Pattern
  -> Option String -> Pattern
```

Key proven properties of substitution (29 theorems, 0 sorries):

- `subst_empty`: empty substitution is the identity;
- `subst_fresh`: substitution on a fresh variable is the identity;
- `commSubst`: the COMM-rule substitution $p[@q/x]$ respects freshness.

7.2 MeTTa Core Interpreter

The `MeTTaCore/` directory provides a complete interpreter specification for Hyperon Experimental MeTTa (3,839 lines, 0 sorries), including:

- `Atom`: the universal term type with `DecidableEq`;
- `Bindings`: variable resolution with merge and transitive lookup;
- `MeTTaState`: the 4-register $\langle i, k, w, o \rangle$ machine;
- `PatternMatch`: bidirectional unification;
- `MinimalOps`: grounded operations $(+, -, *, /, <, \text{etc.})$;
- `RewriteRules`: equation-driven rewriting;

- **Atomspace**: multiset-based knowledge base with query operations;
- **Properties**: progress, confluence, and barbed bisimulation.

8 Type Soundness

Theorem 11 (Substitutability, 0 sorries). *If P and Q are bisimilar processes, then they have the same native types:*

$$P \sim Q \implies \forall \tau. P : \tau \iff Q : \tau$$

Theorem 12 (Type Preservation, 0 sorries). *The COMM and DROP rules preserve types:*

$$\Gamma \vdash P : \tau \wedge P \rightsquigarrow Q \implies \Gamma \vdash Q : \tau$$

Both theorems are proven in `RhoCalculus/Soundness.lean` with a 10-constructor typing judgment `HasType` and explicit typing contexts.

9 ϱ -Calculus Instance

The ϱ -calculus is formalized with full reduction semantics (COMM, DROP), structural congruence (11 rules), and multi-step reduction. The OSLF algorithm derives its spatial-behavioral type system, with the Galois connection $\diamond \dashv \blacksquare$ proven in `RhoCalculus/Soundness.lean` (0 sorries).

10 Type Synthesis: LanguageDef \rightarrow OSLFTypeSystem

The culmination of the formalization is the *type synthesis pipeline*: given any `LanguageDef`, mechanically produce a full `OSLFTypeSystem` with a proven **Galois connection**.

10.1 The Pipeline

The pipeline proceeds in five steps, all implemented in `Framework/TypeSynthesis.lean`:

1. `langReduces lang p q` wraps the executable engine: $q \in \text{rewriteWithContext } lang\ p$.
2. `langRewriteSystem lang` assembles a `RewriteSystem`.
3. `langSpan lang` builds the reduction span (edges = reductions).
4. `langDiamond/langBox` derive modal operators via `derivedDiamond/derivedBox` from the adjoint triple.
5. `langOSLF lang` packages everything into an `OSLFTypeSystem`, with the Galois connection proven by `derived_galois`.

Theorem 13 (Automatic Galois Connection, 0 sorries). *For any LanguageDef lang:*

$$\Diamond_{\text{lang}} \dashv \blacksquare_{\text{lang}}$$

where $\Diamond_{\text{lang}} = \exists_{\text{src}} \circ \text{tgt}^*$ and $\blacksquare_{\text{lang}} = \forall_{\text{tgt}} \circ \text{src}^*$ are derived from the reduction span. No manual proof is needed per language.

11 Language Instances

The pipeline is validated by six language instances of increasing complexity:

11.1 Lambda Calculus (1 sort, 0 crossings)

Untyped lambda calculus with β -reduction (`Framework/LambdaInstance.lean`, 218 lines). One sort (`Term`), two constructors (`App`, `Lam`), one reduction rule. The constructor category is discrete (no sort-crossing constructors). Six executable demos verify β -reduction, multi-step normalization, and formula checking.

11.2 Petri Net (1 sort, 0 crossings)

A simple Petri net with four places and two transitions (`Framework/PetriNetInstance.lean`, 233 lines). Validates that multiset (bag) matching works correctly without any abstraction or substitution machinery. Key properties proven by `native_decide`:

- $\{D\}$ is a dead marking (0 reducts);
- $\{A, B\}$ has exactly 1 reduct via transition T_1 .

11.3 ϱ -Calculus (2 sorts, 2 crossings)

The primary instance with sorts Proc and Name, constructors `NQuote`: `Proc → Name` and `PDrop`: `Name → Proc`. This is the most extensively developed instance, with the full type soundness proof, specialized engine, and Beck–Chevalley analysis (Sections 4–4.5).

11.4 TinyML (2 sorts, 2 crossings, 6 rules)

New: A call-by-value λ -calculus with booleans, pairs, and thunks (`Framework/TinyMLInstance.lean`, 758 lines, 0 sorries).

Listing 10: TinyML language definition (`TinyMLInstance.lean`)

```
-- Sorts: Expr (process sort), Val (data sort)
-- Constructors:
--   Expr: App, If, Fst, Snd, Inject(v:Val)
--   Val: BoolT, BoolF, Lam(`body), PairV, Thunk(e:Expr)
-- Reductions:
--   Beta:     App(Inject(Lam(`body)), Inject(v)) ~> body[v/x]
```

```
-- Force:   Inject(Thunk(e)) ~> e
-- IfTrue:  If(Inject(BoolT), t, e) ~> t
-- IfFalse: If(Inject(BoolF), t, e) ~> e
-- FstPair: Fst(Inject(PairV(a, b))) ~> Inject(a)
-- SndPair: Snd(Inject(PairV(a, b))) ~> Inject(b)
```

TinyML mirrors the ϱ -calculus sort structure:

TinyML	ϱ -Calculus	Role
Expr	Proc	Process sort (carries reduction)
Val	Name	Data sort
Thunk: Expr → Val	NQuote: Proc → Name	Quoting (\diamond)
Inject: Val → Expr	PDrop: Name → Proc	Reflecting (\blacksquare)
Beta	COMM	Main computation rule
Force	DROP	Quoting/reflecting cancel

The CBV strategy is encoded syntactically: β -reduction requires both the function and argument to be wrapped in `Inject(-)`, ensuring arguments are evaluated to values before substitution.

Key proven results:

- `tinyML_crossings`: exactly 2 sort-crossing constructors (`Inject`, `Thunk`)
 - `native_decide`.
- `thunk_is_quoting`: `Thunk` classified as quoting; typing action = \diamond .
- `inject_is_reflecting`: `Inject` classified as reflecting; typing action = \blacksquare .
- `tinyML_typing_action_galois`: $\diamond \dashv \blacksquare$ for TinyML typing actions.
- 11 executable demos including a 3-step reduction chain (`force` → `ifTrue` → `fstPair`).

11.5 MeTTaMinimal (2 sorts, 1 crossing, 2 rules)

A minimal MeTTa state-machine client (`Framework/MeTTaMinimalInstance.lean`, 700 lines, 0 sorries). One sort crossing (`Eval: Expr → State`), 2 rules (rewrite and match-fail). Validates the end-to-end path from checker to Beck-Chevalley graph, including `mettaMinimal_checker_to_BC_graph`: the checker's verdicts are compatible with the categorical semantics.

11.6 MeTTaFull (3 sorts, 2 crossings, 4 rules)

A full MeTTa language definition based on the interpreter specification in `MeTTaCore/ (Framework/MeTTaFullInstance.lean`, 346 lines, 0 sorries). Three sorts (Atom, State, Result), sort-crossing constructors `Interpret: Atom → State` and `Output: State → Result`. Four rewrite rules model the interpret-match-rewrite-output cycle. The end-to-end theorem `mettaFull_checker_to_BC_graph` verifies categorical agreement as with MeTTaMinimal.

12 Verified Formula Checker

The file `Formula.lean` provides a verified bounded model checker for OSLF formulas (582 lines, 0 sorries). The formula type `OSLFFormula` supports \Diamond , \blacksquare , \wedge , \vee , \implies , \top , \perp , and atomic predicates.

Theorem 14 (Checker Soundness, 0 sorries). *If the checker returns `.sat` for formula φ at term p , then $p \models \varphi$ in the denotational semantics.*

The checker supports:

- `checkWithPred`: checking with external predecessor functions, enabling bounded \blacksquare verification;
- `aggregateBox`: universal checking of $\blacksquare\varphi$ over a predecessor list, with proven soundness (`aggregateBox.sat`).

13 Categorical Lift and Presheaf Topos

The file `Framework/CategoryBridge.lean` (1,436 lines, 0 sorries) lifts the Set-level OSLF construction into Mathlib's categorical infrastructure, culminating in the presheaf topos $P(T) = \mathbf{Set}^{T^{\text{op}}}$ over the constructor category.

13.1 Modal Adjunction

The predicate type `Pattern → Prop` carries conflicting category instances (`Preorder.smallCategory` vs. `CategoryTheory.Pi`). We introduce a type wrapper `PredLattice` to disambiguate, then lift:

Listing 11: Modal adjunction (`CategoryBridge.lean`)

```
noncomputable def langModalAdjunction (lang : LanguageDef) :  
  (langGaloisL lang).monotone_l.functor |-  
  (langGaloisL lang).monotone_u.functor :=  
  (langGaloisL lang).adjunction
```

This provides a categorical `Adjunction` between the \Diamond and \blacksquare endofunctors on the predicate preorder category, for any `LanguageDef`.

13.2 Presheaf Topos and Yoneda Representables

Following Williams and Stay [3], we construct the presheaf topos $P(T)$ as a lambda-theory with equality:

Listing 12: Presheaf topos (`CategoryBridge.lean`)

```
def languagePresheafLambdaTheory (lang : LanguageDef) :  
  LambdaTheoryWithEquality where  
  Obj := Functor (Opposite (ConstructorObj lang)) Type  
  -- CCC structure via Mathlib
```

The Yoneda embedding sends each sort s to a representable presheaf $y(s) \in P(T)$. The subobject classifier Ω is constructive via sieves, and the characteristic-map equivalence $\text{Sub}(y(s)) \cong (y(s) \rightarrow \Omega)$ is proven (`languageSortFiber_characteristicEquiv`).

A bridge maps Set-level predicates (`Pattern → Prop`) into presheaf-level subobjects $\text{Sub}(y(s))$ under a naturality condition (`languageSortFiber_ofPatternPred`). The reduction relation is internalized as a `Subfunctor` of the pair presheaf with graph-object diamond/box characterization (`langDiamondUsing_iff_exists_graphStep`, `langBoxUsing_iff_forall_graphIncoming`).

13.3 Current Scope and the Full Grothendieck Construction

Our presheaf topos construction provides the categorical *home* for OSLF types but does not yet fully internalize the modalities. Specifically:

What is proven. $P(T)$ exists as a CCC with finite limits, subobject classifier, and exponentials (all via Mathlib). Yoneda representables, fiber characterization, and the predicate-to-subfunctor bridge are proven. The reduction relation is internal to $P(T)$ as a subfunctor. Beck–Chevalley instances for specific squares (COMM, substitution) are proven at the subobject level.

What remains. The \Diamond/\blacksquare operators are defined as Set-level endomorphisms on (`Pattern → Prop`) and have not yet been lifted to internal natural transformations $\Omega \rightarrow \Omega$ in $P(T)$. This would require:

1. *Internalizing the reduction span*: replacing the constant presheaf with a genuine span $y(\text{Proc}) \leftarrow R \rightarrow y(\text{Proc})$ in $P(T)$, where the edge set varies with the generalized-element context.
2. *Internal existence/universal image*: defining \Diamond/\blacksquare as $\exists_{\text{src}} \circ \text{tgt}^*$ and $\forall_{\text{tgt}} \circ \text{src}^*$ at the subfunctor level.
3. *Agreement*: proving that the internal modalities, evaluated at representable sorts, reproduce the Set-level modalities.

Additionally, a full Grothendieck construction $NT(T) = \int \text{Sub}$ over $P(T)$ (packaging pairs (P, φ) with $P \in P(T)$ and $\varphi \in \text{Sub}(P)$) would complete the connection to Native Type Theory [3]. Mathlib provides `CategoryTheory.Grothendieck` for $\mathcal{C} \rightarrow \mathbf{Cat}$ functors, which could be instantiated with the subobject fibration. Kripke–Joyal forcing semantics, not currently in Mathlib, would be needed to fully connect the internal language of $P(T)$ to OSLF typing judgments.

14 Status and Roadmap

14.1 Milestones

- ✓ **Milestone 1:** Executable ϱ -calculus engine. `reduceStep` computes all one-step reducts; `reduceStep_sound` proven with 0 sorries; 6 executable tests pass.
- ✓ **Milestone 2:** Generic MeTTAIL framework. Language-parametric pattern matching (`Match.lean`) and rewriting (`MeTTAIL/Engine.lean`) for any `LanguageDef`. 8-test agreement suite confirms equivalence with specialized engine.
- ✓ **Milestone 3:** OSLF type synthesis. `langOSLF` mechanically generates an `OSLFTypeSystem` from any `LanguageDef` with auto-proven Galois connection. `GenHasType` provides generated typing rules. Three-layer bridge established.
- ✓ **Milestone 4:** Correctness infrastructure. Totalized `Match.lean` (no `partial def` in core matching). Declarative `DeclReduces` with engine soundness and completeness. Enhanced formula checker with predecessor-based \blacksquare checking.
- ✓ **Milestone 5:** Language instances. Six languages—lambda calculus, Petri nets, ϱ -calculus, TinyML, MeTTaMinimal, and MeTTaFull—each with auto-generated `OSLFTypeSystem` and Galois connection.
- ✓ **Milestone 6:** Presheaf-primary categorical lift. Interface-selected bases (`SortCategoryInterface`) and presheaf-primary predicate fibrations are wired via `CategoryBridge`, including representable-fiber characteristic equivalences and checker-to-fiber soundness.
- ✓ **Milestone 7:** Graph-object BC path and end-to-end clients. `ConstructorCategory`: free category on sort-crossing constructors. `ConstructorFibration`: `SubobjectFibration` + `ChangeOfBase` with proven adjunctions. `ModalEquivalence` and `DerivedTyping`: modal/typing synthesis from constructor structure. `ToposReduction/BeckChevalleyOSLF`: reduction graph objects, explicit substitution squares, and graph-level \Diamond/\blacksquare compatibility consumed by TinyML, MeTTaMinimal, and MeTTaFull checker-to-semantics theorems. All 0 sorries in this core track.

15 Conclusion

We have presented the first machine-checked formalization of the OSLF algorithm, spanning 37,300+ lines of Lean 4 across 111 files with 0 sorries and 0 custom axioms. The formalization demonstrates that the entire pipeline from rewrite rules to spatial-behavioral types can be made rigorous in a modern proof assistant.

The Galois connection at the heart of OSLF is proven at three levels: (1) directly for the ϱ -calculus, (2) categorically for any reduction span via adjoint composition, and (3) lifted to a Mathlib **Adjunction** between endofunctors on the predicate preorder category. All three levels are shown to agree.

The constructor category (Section 4) provides a genuine categorical backbone: sort-crossing constructors become morphisms, the fibered change-of-base gives the adjoint triple $\exists_f \dashv f^* \dashv \forall_f$ for each constructor, and the modal operators \Diamond/\blacksquare are shown to be the typing actions of quoting/reflecting constructors. The derived typing rules (Section 4.4) make explicit that the assignment $\text{NQuote} \mapsto \Diamond$, $\text{PDrop} \mapsto \blacksquare$ is not ad hoc but follows from the constructor’s position in the category. The Beck–Chevalley analysis (Section 4.5) shows that while substitution commutes with change-of-base in the COMM-specific case (proven via `comm_preserves_type`), the universal condition fails—a mathematically interesting negative result proven by explicit counterexample.

The presheaf topos $P(T) = \mathbf{Set}^{T^\text{op}}$ over the constructor category is constructed with Yoneda representables, subobject classifier, and fiber characterization (Section 13). The remaining gap—lifting the Set-level modalities into genuine internal operations of $P(T)$ and instantiating the full Grothendieck construction of Native Type Theory [3]—is precisely characterized in Section 13.3.

The full type synthesis pipeline is validated by six language instances (ϱ -calculus, lambda calculus, Petri nets, TinyML, MeTTaMinimal, and MeTTaFull), each with auto-proven Galois connections. TinyML demonstrates the framework’s ability to handle multi-sort CBV languages with binders, conditionals, and sort-crossing constructors, mirroring the ϱ -calculus’s NQuote/P-Drop structure with Thunk/Inject. The MeTTa instances close the loop from the interpreter specification back to the categorical semantics.

A declarative reduction relation provides an engine-independent specification, and the executable engines are proven both sound and complete with respect to it. The seven *core* implementation milestones are complete and sorry-free, while assumption-audit and literature-alignment items are tracked explicitly as unresolved in `Framework/FULLStatus.lean`.

References

- [1] L. Gregory Meredith and David F. Radestock. A reflective higher-order calculus. *Electronic Notes in Theoretical Computer Science*, 141(5):49–67, 2005.
- [2] L. Gregory Meredith and Mike Stay. Operational semantics in logical form. 2020. Draft / technical report.
- [3] Christian Williams and Mike Stay. Native type theory. In *Proceedings of the 4th Annual International Conference on Applied Category Theory (ACT 2021)*, 2021. arXiv:2211.16865.