

Lean-Backed NS Interface Note

April 21, 2026

Abstract

This note records only the Navier–Stokes material currently backed by Lean in the MeTTape-dia repository. It intentionally omits manuscript-specific claims about Ben Goertzel’s December 11, 2025 SG-Cole-Hopf draft unless they have been encoded and checked in Lean.

The present Lean status is interface-level rather than proof-complete. The formalized material isolates reusable shells for cutoff continuity, topological Cole–Hopf transport, approximation interfaces, finite identity-energy algebra, abstract vorticity bounds, a kernel-semigroup wrapper, and two small concrete models showing both positive approximation instantiations and wrong topology obstructions, including one exact same-state-space split where a summable geometric observable is continuous and a stronger tail detector is not. It also now contains an abstract good observable class, an abstract tail-spike no-go theorem, and a first concrete soft-vs-hard chart fork on that same shared mode space, together with more manuscript-shaped Sobolev-style and energy-style soft-vs-hard forks on that same space, plus a log-energy fork showing that a smoother radial profile still lands on the same good/bad split, and finally a generic energy-profile theorem packaging that split for any continuous radial profile dominated by the identity on nonnegative inputs. It also now contains a theorem-statement-back interface from a local Fefferman-style global-regularity target to the exact uniform vorticity-control tendril currently produced by the grassroots packages, together with a constructive top-down bridge interface decomposing the missing lift into candidate fields, regularity, dynamics, energy, and vorticity-compatibility layers. A new concrete equation-target file now fixes time $\mathbb{R}$, space $\mathbb{R}^3$, and the literal NS equation surface $\partial_t u + (u \cdot \nabla)u + \nabla p = \nu \Delta u$, $\operatorname{div} u = 0$, $u(0, \cdot) = u_0$, and records the honest one-way map from a full vector-valued NS witness to the scalar component shadow currently seen by the Fefferman-style bridge surface. It also now makes one PDE-side slot genuinely concrete: the divergence operator on $\mathbb{R}^3$ is defined as the trace of the spatial Fréchet derivative on the standard basis, and divergence-free initial data are defined by the corresponding vanishing condition. It now also makes the pressure-gradient slot concrete in the same way: the spatial pressure gradient is the actual gradient on $\mathbb{R}^3$, and the momentum equation in the more concrete constructor literally uses that operator. The newest pass also makes the time-derivative slot concrete: $\partial_t u$ is now defined by the actual time Fréchet derivative in the unit time direction, so the current most concrete constructor literally packages $\partial_t u$, ∇p , and $\operatorname{div} u$ on $\mathbb{R} \times \mathbb{R}^3$. It now also makes the convection slot concrete: $(u \cdot \nabla)u$ is defined as the spatial Fréchet derivative applied to the current velocity vector, so the current most concrete constructor now packages the full left-hand side of the NS momentum equation literally, leaving only the Laplacian and analytic regularity slots abstract. It now also makes the Laplacian slot concrete using mathlib’s actual Laplace operator on finite-dimensional inner-product spaces, so the current most concrete constructor packages the full differential operator content of the NS equation on $\mathbb{R} \times \mathbb{R}^3$, leaving only the analytic regularity and energy predicates abstract. The next pass also fixes the smoothness predicates themselves: velocity and pressure are now required to be actual C^∞ maps on $\mathbb{R} \times \mathbb{R}^3$, and initial velocity data are actual C^∞ maps on $\mathbb{R}^3$. The next strengthening then also fixes bounded energy to a literal uniform bound on the spatial kinetic-energy integrals $\int_{\mathbb{R}^3} |u(t, x)|^2 dx$, so what remains abstract is no longer the theorem surface but the actual existence/regularity bridge. The next analytic pass then adds a first concrete continuation surface on top of that explicit theorem

statement: finite-time smooth witnesses on slabs $0 \leq t \leq T$, a literal curl/vorticity operator on $\mathbb{R}^3$, and a uniform-vorticity continuation target saying that such a slab witness should extend to a global one. The next sharpening then adds a closer-to-BKM continuation target using an explicit nonnegative time-envelope $\Omega(t)$ for $\|\omega(t, \cdot)\|$ which is integrable on $[0, T]$. The newest pass also specializes that compatibility frontier back down to the concrete windowed Cole–Hopf heat route: a genuine smooth window still survives the present approximation and heat-envelope shells, and the resulting scaled and affine seed-blend descendants already satisfy bridge-level rigidity theorems for original same-route closure. It does *not* yet formalize the manuscript’s specific chart, local hypotheses $(H1)$, $(H2)$, or the BKM-closing analytic bridge.

1 Scope

The Navier–Stokes Lean work currently lives in:

1. `FluidDynamics/NavierStokes/CutoffContinuity.lean`,
2. `FluidDynamics/NavierStokes/WeakTopologyCountermodel.lean`,
3. `FluidDynamics/NavierStokes/ProductTopologyL1Countermodel.lean`,
4. `FluidDynamics/NavierStokes/ColeHopfTopology.lean`,
5. `FluidDynamics/NavierStokes/ApproximationInterface.lean`,
6. `FluidDynamics/NavierStokes/ChartApproximationModel.lean`,
7. `FluidDynamics/NavierStokes/GeometricModeApproximation.lean`,
8. `FluidDynamics/NavierStokes/ModeStateProductTopologyObstruction.lean`,
9. `FluidDynamics/NavierStokes/ModeStateProductTopologyContinuity.lean`,
10. `FluidDynamics/NavierStokes/ModeStateObservableClass.lean`,
11. `FluidDynamics/NavierStokes/ModeStateTailSensitivityObstruction.lean`,
12. `FluidDynamics/NavierStokes/ModeStateChartFork.lean`,
13. `FluidDynamics/NavierStokes/ModeStateSobolevChartFork.lean`,
14. `FluidDynamics/NavierStokes/ModeStateEnergyChartFork.lean`,
15. `FluidDynamics/NavierStokes/ModeStateLogEnergyChartFork.lean`,
16. `FluidDynamics/NavierStokes/ModeStateEnergyProfileFork.lean`,
17. `FluidDynamics/NavierStokes/FeffermanTargetBack.lean`,
18. `FluidDynamics/NavierStokes/FeffermanTopDownBridge.lean`,
19. `FluidDynamics/NavierStokes/NavierStokesEquationTarget.lean`,
20. `FluidDynamics/NavierStokes/FeffermanCompatibilityFrontier.lean`,
21. `FluidDynamics/NavierStokes/FeffermanWeakSelfModel.lean`,

22. FluidDynamics/NavierStokes/FeffermanSeededSelfModel.lean,
23. FluidDynamics/NavierStokes/FeffermanBoundedSeededSelfModel.lean,
24. FluidDynamics/NavierStokes/FeffermanPressureSeededSelfModel.lean,
25. FluidDynamics/NavierStokes/FeffermanScaledPressureSeededFrontier.lean,
26. FluidDynamics/NavierStokes/FeffermanSeedBlendPressureFrontier.lean,
27. FluidDynamics/NavierStokes/FeffermanSaturatedPressureFrontier.lean,
28. FluidDynamics/NavierStokes/FeffermanFrozenGatePressureFrontier.lean,
29. FluidDynamics/NavierStokes/FeffermanSeedLiveProfileFrontier.lean,
30. FluidDynamics/NavierStokes/FeffermanSeedLiveOperatorFrontier.lean,
31. FluidDynamics/NavierStokes/FeffermanSeedLiveSampleKernelFrontier.lean,
32. FluidDynamics/NavierStokes/FeffermanSeedLiveShiftKernelFrontier.lean,
33. FluidDynamics/NavierStokes/WindowedCutoffApproximation.lean,
34. FluidDynamics/NavierStokes/WindowedColeHopfHeatPackage.lean,
35. FluidDynamics/NavierStokes/WindowedColeHopfHeatFoundationFrontier.lean,
36. FluidDynamics/NavierStokes/WindowedColeHopfHeatSeededModel.lean,
37. FluidDynamics/NavierStokes/WindowedColeHopfHeatScaledSeededFrontier.lean,
38. FluidDynamics/NavierStokes/WindowedColeHopfHeatSelfCompatibilityObstruction.lean,
39. FluidDynamics/NavierStokes/WindowedColeHopfHeatSeedBlendFrontier.lean,
40. FluidDynamics/NavierStokes/WindowedColeHopfHeatFrozenGateFrontier.lean,
41. FluidDynamics/NavierStokes/WindowedColeHopfHeatFrozenGateObstruction.lean,
42. FluidDynamics/NavierStokes/WindowedColeHopfHeatSeedLiveProfileFrontier.lean,
43. FluidDynamics/NavierStokes/WindowedColeHopfHeatSeedLiveProfileObstruction.lean,
44. FluidDynamics/NavierStokes/WindowedColeHopfHeatSeedLiveOperatorFrontier.lean,
45. FluidDynamics/NavierStokes/WindowedColeHopfHeatSeedLiveOperatorObstruction.lean,
46. FluidDynamics/NavierStokes/WindowedColeHopfHeatSeedOnlyOperatorObstruction.lean,
47. FluidDynamics/NavierStokes/WindowedColeHopfHeatLiveOnlySampleKernelObstruction.lean,
48. FluidDynamics/NavierStokes/WindowedColeHopfHeatSampleKernelFrontier.lean,
49. FluidDynamics/NavierStokes/WindowedColeHopfHeatShiftKernelFrontier.lean,
50. FluidDynamics/NavierStokes/WindowedColeHopfHeatSeedBlendObstruction.lean,

51. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSampleKernelObstruction.lean`,
52. `FluidDynamics/NavierStokes/WindowedColeHopfHeatShiftKernelObstruction.lean`,
53. `FluidDynamics/NavierStokes/WindowedColeHopfHeatMaskedShiftKernelObstruction.lean`,
54. `FluidDynamics/NavierStokes/WindowedColeHopfHeatGeneralShiftKernelObstruction.lean`,
55. `FluidDynamics/NavierStokes/WindowedColeHopfHeatGeneralShiftKernelMaskedSpecialization.lean`,
56. `FluidDynamics/NavierStokes/WindowedColeHopfHeatGeneralShiftKernelIdentitySampleSpecialization.lean`,
57. `FluidDynamics/NavierStokes/WindowedColeHopfHeatIdentitySampleKernelObstruction.lean`,
58. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedFrontier.lean`,
59. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedObstruction.lean`,
60. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedSampleKernelObstruction.lean`,
61. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedMaskedShiftKernelObstruction.lean`,
62. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedShiftKernelObstruction.lean`,
63. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedShiftKernelConstantMagnitudeFrontier.lean`,
64. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedShiftKernelInvariantFrontier.lean`,
65. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedShiftKernelInvariantObstruction.lean`,
66. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedAnchoredShiftInvariantFrontier.lean`,
67. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedAnchoredShiftInvariantObstruction.lean`,
68. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedShiftKernelInvariantAnchoredSpecialization.lean`,
69. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedShiftKernelMaskedSpecialization.lean`,
70. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedShiftKernelIdentitySampleSpecialization.lean`,
71. `FluidDynamics/NavierStokes/WindowedColeHopfHeatSaturatedIdentitySampleKernelObstruction.lean`,
72. `FluidDynamics/NavierStokes/IdentityEnergy.lean`,
73. `FluidDynamics/NavierStokes/ColeHopfInterface.lean`,
74. `FluidDynamics/NavierStokes/KernelSemigroupInterface.lean`.

These modules form the live NS interface shell. The currently verified slice is recorded explicitly in Section 5. Together they provide a formal shell for what a future SG-Cole-Hopf proof attempt would need to instantiate.

2 Lean-Backed Results

2.1 Cutoff Continuity Shell

In `CutoffContinuity.lean`, the generic cutoff-shaped potential

$$\text{cutoffPotential}(x) = \text{cutoff}(\text{radius}(x)) \cdot \text{statistic}(x)$$

is formalized together with:

1. continuity of the cutoff potential when the scalar pieces are continuous;
2. a sequential obstruction theorem: if the cutoff factor and statistic converge to the wrong product, then the cutoff potential cannot converge to its claimed value;
3. the corresponding pointwise non-continuity criterion.

So Lean already captures the generic shape of a chart-cutoff continuity argument and its failure mode.

2.2 Concrete Weak-Topology Countermodel

In `WeakTopologyCountermodel.lean`, the abstract cutoff/Cole–Hopf shell is instantiated on the concrete state space $\mathbb{R} \times \mathbb{R}$ equipped with the topology induced by the first coordinate. Lean proves:

1. a sequence can converge in the ambient topology while changing a hidden coordinate;
2. a cutoff potential depending on that hidden coordinate is not continuous;
3. the associated Cole–Hopf transform is likewise not continuous.

This is a formal countermodel for the generic wrong-topology failure mode: ambient convergence can be too weak for the cutoff datum it is supposed to control.

2.3 Infinite-Mode Product-Topology Countermodel

In `ProductTopologyL1Countermodel.lean`, Lean moves the same mismatch to the raw infinite-mode sequence space $\mathbb{N} \rightarrow \mathbb{R}$ with the product topology. It proves:

1. finite-mode truncations converge coordinatewise in the product topology;
2. a tail spike moving to higher and higher coordinates converges to the zero sequence in that topology;
3. the stronger tail-sensitive observable $\sum_n |x_n|$ stays equal to 1 on that spike sequence and is therefore not continuous at zero;
4. the associated Cole–Hopf datum is likewise not continuous.

This is the first Lean-backed infinite-mode obstruction in the NS lane that directly separates weak coordinatewise convergence from a stronger tail norm.

2.4 Topological Cole–Hopf Shell

In `ColeHopfTopology.lean`, Lean formalizes:

1. continuity of the scalar map $y \mapsto \exp(-y/(2\nu))$;
2. continuity of the induced field-level Cole–Hopf transform;
3. injectivity when $\nu \neq 0$;
4. transfer of potential-limit mismatch into Cole–Hopf-limit mismatch;
5. the same obstruction specialized to cutoff-shaped potentials.

This means the topological $W \mapsto \phi$ step is already packaged in Lean, provided one supplies continuity or convergence of the underlying potential.

2.5 Approximation Interface

In `ApproximationInterface.lean`, Lean packages an abstract positive repair obligation:

if an approximation scheme makes the chart radius observable and the chart statistic converge, then the cutoff potential and the Cole–Hopf datum also converge.

This is useful because it isolates exactly what an approximation/truncation theorem would have to prove in order for the cutoff/Cole–Hopf layer to behave well.

2.6 Concrete Chart/Truncation Model

In `ChartApproximationModel.lean`, the abstract approximation interface is instantiated on the concrete chart state space $\mathbb{R} \times \mathbb{R}$ with:

1. a hidden-coordinate radius observable;
2. a simple mixed chart statistic;
3. a truncation map that drops the hidden coordinate only at stage 0 and is exact thereafter.

Lean then proves:

1. convergence of the radius observable under truncation;
2. convergence of the chart statistic under truncation;
3. the resulting convergence of the cutoff potential;
4. the resulting convergence of the induced Cole–Hopf datum.

So the approximation interface is now backed by a concrete positive model, not just an abstract shell.

2.7 Geometric Finite-Mode Approximation Model

In `GeometricModeApproximation.lean`, the same approximation interface is instantiated on bounded infinite coefficient sequences. Lean formalizes:

1. a geometrically weighted mode radius;
2. a chart statistic built from the zeroth visible mode plus that radius;
3. finite-mode truncation keeping only the first N coefficients;
4. summability of the weighted mode series under the uniform coefficient bound;
5. convergence of the weighted radius under finite-mode truncation;
6. convergence of the derived chart statistic, the cutoff potential, and the induced Cole–Hopf datum.

This is the first genuinely infinite-mode positive instantiation in the NS Lean lane: the approximation is not just “wrong at stage 0, exact thereafter,” but a real truncation-to-limit model on an infinite sequence state space.

2.8 Same-State-Space Product-Topology Obstruction

In `ModeStateProductTopologyObstruction.lean`, Lean returns to the exact `ModeState` and `truncateModes` objects from `GeometricModeApproximation.lean`, but equips that same state space with the coordinatewise product topology induced by the coefficient map. It proves:

1. the existing finite-mode truncations `truncateModes` do converge to the target state in this product topology;
2. a spike sequence still converges to the zero state in that same topology;
3. a stronger tail-presence observable, which records whether any nonzero coefficient exists, is not continuous there;
4. the associated Cole–Hopf datum is likewise not continuous.

So Lean now contains a positive and a negative result on one common infinite-mode object: the geometric weighted radius behaves well under truncation, while a stronger tail-sensitive observable on the same state space does not.

2.9 Same-State-Space Product-Topology Continuity

In `ModeStateProductTopologyContinuity.lean`, Lean proves the matching positive theorem on the exact same `ModeState` equipped with the exact same coordinatewise product topology:

1. each coordinate observable $x \mapsto x_n$ is continuous;
2. each weighted mode term $x \mapsto |x_n|2^{-n}$ is continuous;
3. the summable geometric radius $\sum_n |x_n|2^{-n}$ is continuous;
4. the derived chart statistic, cutoff potential, and Cole–Hopf datum are therefore continuous as well.

So the current NS Lean lane now cleanly separates two observables on one common infinite-mode state space: a summably dominated geometric radius that behaves well in the product topology, and a stronger tail detector that does not.

2.10 Abstract Good Observable Class

In `ModeStateObservableClass.lean`, Lean packages the good side on the shared `ModeState`: nonnegative summable mode weights. For that whole class, Lean proves:

1. continuity of the weighted radius observable;
2. continuity of the derived statistic, cutoff potential, and Cole–Hopf datum;
3. convergence of all of these along the existing truncations `truncateModes`;
4. recovery of the earlier geometric radius model as a special case.

So the good side is no longer just one geometric example. It is now an abstract summable observable class on the shared mode space.

2.11 Abstract Tail-Spike No-Go Theorem

In `ModeStateTailSensitivityObstruction.lean`, Lean packages the bad side abstractly too. If an observable is constant on every remote single-spike state `tailSpikeMode(N)` but takes a different value at `modeZero`, Lean proves:

1. the raw observable is not continuous at `modeZero`;
2. the same obstruction lifts to cutoff potentials;
3. the same obstruction lifts further to Cole–Hopf data.

So the bad side is no longer tied only to the one hand-built `tailPresenceRadius`. It is now a reusable theorem pattern.

2.12 First Concrete Chart Fork

In `ModeStateChartFork.lean`, Lean forces the first explicit chart-style fork on the shared mode space:

1. the soft support observable $t \mapsto |t|/(1 + |t|)$, combined with summable mode weights, lies on the good side and yields a continuous truncation-friendly chart radius, statistic, cutoff potential, and Cole–Hopf datum;
2. the hard support observable, which only asks whether any coefficient is nonzero, lies on the bad side and fails continuity at `modeZero`, together with its cutoff and Cole–Hopf lifts.

This is the first genuine theorem-level NS fork in Lean:

soft support belongs to the good summable side, hard support belongs to the bad tail-sensitive side.

2.13 Sobolev-Style Chart Fork

In `ModeStateSobolevChartFork.lean`, Lean pushes that fork one step closer to manuscript shape by amplifying each mode by the polynomial factor $((n + 1)^m)$ before applying the chart cutoff:

1. the soft Sobolev observable, built from the bounded scalar $t \mapsto |t|/(1 + |t|)$ together with a concrete inverse-square envelope, is continuous and truncation-friendly;
2. the corresponding hard Sobolev observable, which only asks whether any amplified mode is nonzero, is proved equal to the earlier hard support detector and is therefore not continuous at `modeZero`, together with its cutoff and Cole–Hopf lifts.

So the current Lean fork is now sharper:

even after Sobolev-style high-frequency amplification, the soft bounded version lands on the good summable side, while the hard support version still lands on the bad side.

2.14 Energy-Style Chart Fork

In `ModeStateEnergyChartFork.lean`, Lean pushes the same shared-state fork one step further toward an energy-style observable:

1. the soft energy observable squares the Sobolev-amplified modes and then applies a compensating inverse-power envelope; Lean proves that this radius is continuous and truncation-friendly, and so are its cutoff and Cole–Hopf lifts;
2. the hard energy observable asks whether any amplified squared mode crosses the threshold 1; Lean proves that this detector is stable on the remote spike family, disagrees at `modeZero`, and is therefore not continuous, again together with its cutoff and Cole–Hopf lifts.

So the current shared-mode-space fork is now sharper:

if the chart observable keeps enough inverse-power compensation to remain summable after Sobolev amplification and squaring, Lean puts it on the good approximation-friendly side; if it turns into a hard energy threshold on those amplified modes, Lean puts it on the bad tail-sensitive side.

2.15 Log-Energy Chart Fork

In `ModeStateLogEnergyChartFork.lean`, Lean tests a smoother radial profile on the same shared mode space by replacing raw squared Sobolev energy by $\log(1 + \text{energy})$:

1. the soft log-energy observable is continuous and truncation-friendly; the key formal domination is $\log(1 + u) \leq u$, which lets Lean compare the log-energy terms to the already-controlled soft energy terms;
2. the hard-side analogue asking whether $\log(1 + \text{energy}) > 0$ still lands on the bad side: it is stable on the remote spike family, disagrees at `modeZero`, and therefore its cutoff and Cole–Hopf lifts are also not continuous.

So the shared-mode-space fork is now sharper again:

even replacing raw energy by the smoother profile $\log(1 + \text{energy})$ does not escape the same dichotomy: the summably compensated soft version stays on the good side, while the hard detection version stays on the bad side.

2.16 Generic Energy-Profile Fork

In `ModeStateEnergyProfileFork.lean`, Lean abstracts the preceding examples into one theorem pattern. An `EnergyProfile` is a globally continuous radial profile $\rho : \mathbb{R} \rightarrow \mathbb{R}$ satisfying, on nonnegative inputs:

1. $0 \leq \rho(u)$,
2. $\rho(u) \leq u$,
3. $\rho(0) = 0$,
4. $u > 0 \Rightarrow \rho(u) > 0$.

For every such profile, Lean proves:

1. the summably compensated soft radius is continuous and truncation-friendly;
2. the induced hard detector agrees with the raw hard Sobolev support detector;
3. that hard detector is tail-spike stable and therefore not continuous at `modeZero`, together with its cutoff and Cole–Hopf lifts.

The file also includes canonical instances:

1. the identity profile, recovering the raw soft energy radius;
2. the clipped log profile $u \mapsto \log(1 + \max(u, 0))$, recovering the soft and hard log-energy constructions on the nonnegative energy inputs that actually occur;
3. a bounded rational saturation profile $u \mapsto \max(u, 0)/(1 + \max(u, 0))$, which is closer to a chart-cutoff saturation shape and whose hard detector Lean proves collapses back to the hard Sobolev support detector;
4. a smooth exponential saturation profile $u \mapsto 1 - \exp(-\max(u, 0))$, again with the same hard-detector collapse;
5. a smooth-transition profile $u \mapsto u \cdot \text{smoothTransition}(u)$, using `mathlib's Real.smoothTransition`, again with the same hard-detector collapse;
6. more strongly, a positive-parameter family $u \mapsto u \cdot \text{smoothTransition}(cu)$ for $c > 0$, all with the same hard-detector collapse;
7. stronger still, an affine two-parameter family $u \mapsto u \cdot \text{smoothTransition}(cu + d)$ for $c > 0, d \geq 0$, all with the same hard-detector collapse.

The new boundary result is that a genuinely windowed difference-of-transitions profile already leaves the shell. Lean defines

$$u \mapsto u \cdot (\text{smoothTransition}(cu + 1) - \text{smoothTransition}(cu))$$

and proves that for every $c > 0$ this profile cannot be the `profile`-field of any `EnergyProfile`: it vanishes at the positive input $u = c^{-1}$, so it cannot satisfy the required strict positivity on all positive energies.

So the current shared-mode-space NS fork is now theorem-level rather than example-level:

every continuous radial profile that stays nonnegative, vanishes only at zero, and is dominated by the identity on nonnegative inputs lands on the same good/bad fork.

2.17 Finite Identity-Energy Algebra

In `IdentityEnergy.lean`, Lean formalizes a finite carré-du-champ-style energy

$$\gamma(D) = \sum_i D_i^2$$

and proves:

1. basic positivity and scaling laws for γ ;
2. the log-gradient comparison used in the identity-only Cole–Hopf calculation;
3. the finite Cauchy–Schwarz frame-evaluation bound;
4. the resulting pointwise vorticity-style estimate after push-down through a curl frame.

So the finite algebraic spine of the identity-only bound is already Lean-backed.

2.18 Abstract Cole–Hopf/Vorticity Interface

In `ColeHopfInterface.lean`, the finite identity-energy algebra is packaged into an abstract structure carrying:

1. uniform positivity of $\Phi(t)$;
2. a uniform identity-energy bound on its directional derivatives;
3. a uniform frame-curl bound.

From those hypotheses, Lean derives:

1. a bound on the log-potential coefficients $W(t) = -2\nu \log \Phi(t)$;
2. a pointwise abstract vorticity bound;
3. the uniform-in-time version of the same estimate.

2.19 Kernel-Semigroup Wrapper

In `KernelSemigroupInterface.lean`, the same identity-only Cole–Hopf/vorticity algebra is wrapped around a local Markov-kernel semigroup. This does not add new analytic content. It says that if a future NS evolution is encoded as a local kernel semigroup, then the existing identity-energy bridge can be reused unchanged.

2.20 Fefferman-Style Target Backward Interface

In `FeffermanTargetBack.lean`, Lean now works backward from a local mirror of the Clay/Fefferman output clause. The file first defines a `FeffermanPredicateKit`, so the target is parameterized by abstract output-side predicates but still requires actual proofs of them. It then does three things:

1. it defines `UniformVorticityTendril`, the exact smallest theorem-statement-back fragment currently reached by the grassroots lane: a globally defined vorticity field with one uniform envelope;

2. it proves that the existing local shells `ColeHopfIdentityData`, `FiniteModeColeHopfData`, `ColeHopfKernelSemigroupData`, `FiniteModeColeHopfKernelData`, `ManuscriptTruncationPackage`, and `SharedApproximationPackage` all land in that same tendril;
3. it defines `FeffermanLiftObligations`, the exact additional lift still missing if one wants to climb from that tendril to a local Fefferman-style global-regularity output witness: velocity, pressure, proofs of the chosen smoothness / dynamics / energy predicates, and a proof of the chosen vorticity-to-velocity compatibility predicate.

So the target-back picture is now explicit in Lean:

current grassroots NS packages already produce a uniform vorticity-control tendril, but a separate lift is still required before one has even a local mirror of the Fefferman-style global-regularity output clause.

2.21 Top-Down Fefferman Bridge

In `FeffermanTopDownBridge.lean`, Lean turns that missing lift into a positive program rather than a vague gap. Relative to a chosen predicate kit and a chosen vorticity-compatibility predicate, the file defines:

1. `VelocityPressureCandidate`, the candidate output fields;
2. `RegularityCertificate`, `DynamicsCertificate`, and `EnergyCertificate`, separating smoothness, PDE/incompressibility / initial-condition, and bounded-energy obligations;
3. `VorticityCompatibility`, isolating the extra theorem that ties the already-available vorticity tendril back to the chosen velocity field;
4. `TopDownFeffermanBridge`, bundling exactly those layers.

Lean then proves that any such bridge assembles:

1. the explicit `FeffermanLiftObligations` from the target-back file;
2. the local Fefferman-style global-regularity clause;
3. preservation of the original uniform vorticity tendril.

Finally, the file proves package-level realization theorems for the current NS lane: if a top-down bridge is supplied on the tendril induced by `ColeHopfIdentityData`, `FiniteModeColeHopfData`, `ColeHopfKernelSemigroupData`, `FiniteModeColeHopfKernelData`, `ManuscriptTruncationPackage`, or `SharedApproximationPackage`, then the local Fefferman-style clause follows immediately.

So the constructive NS picture is now also explicit in Lean:

the current grassroots route already reaches a uniform vorticity tendril, and the remaining top-down work is decomposed into theorem-sized lift layers rather than one monolithic miracle step.

2.22 Concrete Navier–Stokes Equation Target

In `NavierStokesEquationTarget.lean`, Lean now states the first literal NS theorem surface in the current lane. The file fixes time to $\mathbb{R}$, space to $\mathbb{R}^3$, velocity to a vector field $u : \mathbb{R} \rightarrow \mathbb{R}^3 \rightarrow \mathbb{R}^3$, and pressure to a scalar field $p : \mathbb{R} \rightarrow \mathbb{R}^3 \rightarrow \mathbb{R}$. Relative to a chosen differential-operator kit, it defines:

1. the pointwise momentum equation

$$\partial_t u + (u \cdot \nabla)u + \nabla p = \nu \Delta u;$$

2. the incompressibility condition $\operatorname{div} u = 0$;
3. the initial condition $u(0, \cdot) = u_0$;
4. a concrete global-regularity witness and the corresponding existence clause for smooth divergence-free initial data with positive viscosity.

The same file also records the honest interface to the current scalar Fefferman-style bridge surface: every full vector-valued NS witness induces, for each coordinate $i \in \{0, 1, 2\}$, a scalar component witness for an associated `FeffermanPredicateKit`. So the current route is now pinned against an actual PDE-side target, but only through its scalar component shadow.

The next small but real strengthening is that one operator is no longer merely named. In the partially concrete constructor `mkWithConcreteDivergence`, the divergence slot is fixed to the trace of the spatial Fréchet derivative on the standard basis of $\mathbb{R}^3$, and Lean proves that incompressibility and divergence-free initial data are exactly the vanishing of those explicit divergence expressions.

The next strengthening fixes a second PDE-side slot. In the more concrete constructor `mkWithConcreteGradientAndDivergence`, the pressure-gradient slot is the actual spatial gradient on $\mathbb{R}^3$, so the momentum equation is no longer merely named there: Lean proves that it is literally the equation with ∇p given by that explicit gradient operator.

The next strengthening fixes a third PDE-side slot. In `mkWithConcreteTimeGradientAndDivergence`, the time-derivative slot is the actual time Fréchet derivative in the unit time direction, so the momentum equation is now literally the equation with explicit $\partial_t u$, ∇p , and $\operatorname{div} u$ on $\mathbb{R} \times \mathbb{R}^3$, while divergence-free initial data remain the explicit vanishing condition from the concrete divergence constructor.

The next strengthening fixes a fourth PDE-side slot. In `mkWithConcreteTimeConvectionGradientAndDivergence`, the convection term $(u \cdot \nabla)u$ is the spatial Fréchet derivative applied to the current velocity vector, so the momentum equation is now literally the equation with explicit $\partial_t u$, $(u \cdot \nabla)u$, ∇p , and $\operatorname{div} u$ on $\mathbb{R} \times \mathbb{R}^3$; only the Laplacian and the analytic regularity/energy predicates remain abstract at this surface.

The next strengthening fixes the fifth PDE-side slot. In `mkWithConcreteDifferentialOperators`, the Laplacian is the actual mathlib Laplace operator on the fixed-time velocity slice, so the momentum equation is now literally the full differential equation with explicit $\partial_t u$, $(u \cdot \nabla)u$, ∇p , Δu , and $\operatorname{div} u$ on $\mathbb{R} \times \mathbb{R}^3$. At that surface, what remains abstract is no longer the differential operator content of the PDE, but the analytic regularity, energy, and existence machinery.

The next strengthening fixes the smoothness predicates themselves. In `mkWithConcreteDifferentialOperators`, velocity and pressure are required to be literal C^∞ maps on $\mathbb{R} \times \mathbb{R}^3$, while initial velocity data are literal C^∞ maps on $\mathbb{R}^3$. So at this surface the remaining abstract slot is no longer “smoothness” in name only: it is essentially the bounded-energy predicate and the actual existence/regularity theorem. An important local repair was needed here after mathlib’s current `ContDiff` indexing semantics became sharper: `ContDiff` $\mathbb{R} \top$ is an ω -level analytic requirement, not the intended C^∞

surface. The local NS theorem surface now explicitly uses `ContDiff ℝ ∞` for velocity, pressure, and initial data, so the formalized smoothness slot once again matches the manuscript-level “smooth” / Schwartz reading rather than an unintended analyticity demand.

The next strengthening fixes that remaining theorem-surface predicate too. In `mkFullyConcreteNavierStokesS` the bounded-energy slot is literally the statement that there exists a uniform constant $C \geq 0$ such that for every time t , the kinetic-energy density $|u(t, x)|^2$ is integrable on $\mathbb{R}^3$ and

$$\int_{\mathbb{R}^3} |u(t, x)|^2 dx \leq C.$$

So the local target is now fully concrete at the PDE-surface level: differential operators, smoothness, incompressibility, initial condition, and bounded energy are all literal Lean definitions on $\mathbb{R} \times \mathbb{R}^3$. What remains abstract is the actual theorem asserting global existence and regularity for that fully concrete surface.

The next tightening removes one more layer of packaging. The same file now defines an explicit proposition saying directly that for every viscosity $\nu > 0$ and every smooth divergence-free initial velocity u_0 , there exist u and p such that u and p are smooth,

$$\partial_t u + (u \cdot \nabla)u + \nabla p = \nu \Delta u, \quad \operatorname{div} u = 0, \quad u(0, \cdot) = u_0,$$

and the kinetic-energy integrals are uniformly bounded in time. Lean proves that this fully explicit statement is equivalent to the current structured `ConcreteNavierStokesMillenniumTarget`. So the target is no longer only concrete “up to record fields”; it is now pinned to a literal theorem statement.

The first bottom-up follow-up file, `VectorCalculusR3.lean`, extracts a reusable concrete vector-calculus layer from that theorem surface. It defines coordinate derivatives and the literal curl/vorticity operator on $\mathbb{R}^3$, then proves real linearity and zero lemmas for the concrete operators ∂_t , div , ∇p , and Δ , plus the corresponding coordinate-derivative facts and concrete vorticity identities (zero field and scalar-multiplication commutation). This is meant as curriculum-style material with independent value, not more interface packaging.

The next file, `NavierStokesUniformVorticityContinuationTarget.lean`, adds the first concrete analytic bridge surface on top of that literal theorem. It defines explicit finite-time smooth witnesses on a slab $0 \leq t \leq T$, the concrete vorticity operator $\nabla \times u$ on $\mathbb{R}^3$, a uniform vorticity bound on the slab, and the corresponding continuation target saying that such a witness should extend to a global smooth incompressible solution with bounded kinetic energy. Lean also proves the sanity theorem that the fully explicit global theorem target subsumes this continuation target, and hence so does the structured fully concrete target via the explicit-equivalence theorem above. At present this is only a one-way subsumption result: the Lean proof simply reuses the full global theorem target and does not yet use the vorticity-bound hypothesis in any analytic argument. But this file now also contains an operational zero-data instantiation: it constructs a canonical finite-time zero witness with uniform zero vorticity bound and proves that, if the uniform continuation target is assumed, one gets an explicit global output for zero initial data at any positive viscosity. In addition, this same file now proves a target-free baseline theorem: `ExplicitConcreteNavierStokesGlobalOutput_zero`, together with the matching clause-level inhabitant `ExplicitConcreteNavierStokesRegularityClause_zero`. Those two lemmas do not assume any continuation target; they are direct concrete PDE-side sanity anchors on the literal $\mathbb{R} \times \mathbb{R}^3$ statement surface. The same concrete layer now also exposes a real pressure-gauge symmetry: `smoothSpaceTimePressure_contDiff_spatialSlice` and `smoothSpaceTimePressure_differentiableAt_spatialSlice` give the fixed-time smoothness

bridge for pressure slices, `smoothSpaceTimeVelocity_add`, `smoothSpaceTimeVelocity_sub`, and `smoothSpaceTimeVelocity_linearCombination` now make the corresponding smooth velocity algebra explicit on $\mathbb{R} \times \mathbb{R}^3$, `smoothSpaceTimePressure_const`, `smoothSpaceTimePressure_const_smul`, `smoothSpaceTimePressure_sub`, and `smoothSpaceTimePressure_linearCombination` do the same for smooth pressure fields, `spatialPressureGradient_add`, `spatialPressureGradient_const_smul`, `spatialPressureGradient_sub`, and `spatialPressureGradient_linearCombination` give the matching affine algebra for concrete pressure gradients on each fixed spatial slice, `ExplicitFiniteTimeRegularityWitness.addPressureOfZeroSpatialGradient` and `ExplicitConcreteNavierStokesGlobalOutput.addPressureOfZeroSpatialGradient` now package the gauge transport principle itself at the explicit witness and global-output layers, while `ExplicitFiniteTimeRegularityWitness.uniformVorticityBound_addPressureOfZeroSpatialGradient` and `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_addPressureOfZeroSpatialGradient` push the same invariance into the two concrete continuation-data packages, `smoothSpaceTimePressure_timeOnly` records that any smooth scalar function of time alone gives a smooth space-time pressure field, and `spatialPressureGradient_timeOnly` records that such a time-only pressure has zero spatial gradient. The resulting witness/global-output transport theorems `ExplicitFiniteTimeRegularityWitness.addTimeOnlyPressure` and `ExplicitConcreteNavierStokesGlobalOutput.addTimeOnlyPressure` show that adding any smooth pressure gauge of the form $p(t)$ leaves the concrete Navier–Stokes data valid. The older constant-pressure transport theorems are now honest special cases of this larger gauge symmetry. Consequently Lean now also proves the explicit time-gauge baseline `ExplicitConcreteNavierStokesGlobalOutput_zero_addPressureOfZeroSpatialGradient` and `ExplicitFiniteTimeRegularityWitness.addTimeOnlyPressure` zero velocity with arbitrary smooth time-only pressure is still a genuine global solution on the literal PDE surface. The same baseline is now packaged at the regularity-clause level too by `ExplicitConcreteNavierStokesRegularityClause_addPressureOfZeroSpatialGradient`, `ExplicitFiniteTimeRegularityWitness.addTimeOnlyPressure`, `ExplicitConcreteNavierStokesGlobalOutput_zero_addPressureOfZeroSpatialGradient`, `ExplicitConcreteNavierStokesRegularityClause_zero_addPressureOfZeroSpatialGradient`, `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_zero_addPressureOfZeroSpatialGradient` which package the generic slice-wise zero-gradient gauge symmetry first on the arbitrary-data explicit clause layer and then on the zero baseline, and the older constant-pressure and time-only baseline lemmas are now just wrappers around those generic transports: `ExplicitConcreteNavierStokesRegularityClause_zero_addPressureOfZeroSpatialGradient`, `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_zero_constantPressure`, `ExplicitConcreteNavierStokesGlobalOutput_zero_addPressureOfZeroSpatialGradient` and `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_zero_timeOnlyPressure`. The same generic transport now also reaches the structured fully concrete whole-space clause layer via `concreteNavierStokesGlobalRegularityClause_addPressureOfZeroSpatialGradient` and `finiteEnergyConcreteNavierStokesGlobalRegularityClause_addPressureOfZeroSpatialGradient`, with the zero-data instances exposed via `concreteNavierStokesGlobalRegularityClause_zero_addPressureOfZeroSpatialGradient` and `finiteEnergyConcreteNavierStokesGlobalRegularityClause_zero_addPressureOfZeroSpatialGradient` so the concrete clause layer itself inherits this pressure-gauge symmetry. This same pressure-gauge invariance now also reaches the uniform continuation side: `ExplicitUniformVorticityContinuationClause_apply_of_uniformVorticityBound_addTimeOnlyPressure` shows that a same-horizon uniform continuation clause can be applied after a smooth time-only pressure-gauge change, while `ExplicitFiniteTimeRegularityWitness.uniformVorticityBound_addTimeOnlyPressure` shows that the same uniform bound survives gauge change plus restriction to a shorter slab, and `ExplicitUniformVorticityContinuationClause_apply_of_uniformVorticityBound_addTimeOnlyPressure_restrict_has_uniformVorticityBound` turns that directly into a smaller-horizon operational clause theorem. Lean now also instantiates that operational theorem concretely at zero data first for arbitrary slice-wise zero-gradient gauges: `zeroFiniteTimeRegularityWitness_addPressureOfZeroSpatialGradient_restrict_has_uniformVorticityBound` shows that the gauged-and-restricted zero witness itself already carries the explicit shorter-slab zero vorticity bound, and `ExplicitUniformVorticityContinuationClause_implies_zeroGlobalOutput_addPressureOfZeroSpatialGradient` shows that a smaller-slab uniform continuation clause can therefore be applied directly to that concrete gauged zero witness. The older `..._timeOnlyPressure` and constant-pressure lemmas are now just wrappers around that general zero-spatial-gradient transport. The same zero-data operational

packaging now exists on the BKM side too via `zeroFiniteTimeRegularityWitness_addPressureOfZeroSpatialG`, `ExplicitBKMContinuationClause_implies_zeroGlobalOutput_addPressureOfZeroSpatialGradient`, and their repaired finite-energy analogues. The same layer now also includes the direct clause-level formulation `ExplicitUniformVorticityContinuationClause_implies_zeroGlobalOutput`, which applies one concrete uniform continuation clause instance to the canonical zero slab witness. The premise is now also packaged explicitly by `zeroFiniteTimeRegularityWitness_exhibits_uniformContinuation` which says the quantified witness-and-bound hypothesis is genuinely inhabited at zero data and therefore not merely vacuous there. The zero-data route is now factored through a reusable uniform bound lemma `uniformVorticityBoundUpTo_zero`. This layer now also contains a reusable rescaling transport lemma `uniformVorticityBoundUpTo_const_smul`: if $\|\omega_u\| \leq B$ on a slab, then $\|\omega_{a \cdot u}\| \leq |a| B$ on the same slab, plus monotonicity and sign-flip transport lemmas `uniformVorticityBoundUpTo_mono` and `uniformVorticityBoundUpTo_neg`. The concrete vector-calculus layer underneath now also exposes `smoothSpaceTimeVelocity_contDiff_spatialSlice` and `smoothSpaceTimeVelocity_differentiableAt` which turn full space-time smoothness into fixed-time spatial differentiability. Using that bridge, the same file now proves the additive closure theorem `uniformVorticityBoundUpTo_add`: smooth slab bounds for u and v combine into a slab bound for $u + v$ with parameter $B + B'$. So the uniform layer now supports not only rescaling and restriction, but also explicit composition of concrete vorticity bounds under field addition. It also has time-restriction transport lemmas `uniformVorticityBoundUpTo_restrict` and `boundedKineticEnergyUpTo_restrict`. It also now adds a witness-level restriction constructor `ExplicitFiniteTimeRegularityWitness_restrict` and the induced transport lemma `ExplicitFiniteTimeRegularityWitness.uniformVorticityBound_restrict`, as well as the clause-level horizon-monotonicity theorem `ExplicitUniformVorticityContinuationClause_mono_h`. The same bottom-up lane now also contains an explicit nonzero-looking sanity family that exposes a genuine obstruction instead of more interface layering. The vector-calculus file introduces `constantVelocityField` and `constantInitialVelocity` and proves that such fields have zero time derivative, zero convection, zero Laplacian, zero divergence, and zero vorticity, so they are honest differential-side constant-field solutions on the literal $\mathbb{R} \times \mathbb{R}^3$ PDE surface. The uniform-vorticity file then proves `uniformVorticityBoundUpTo_constantVelocityField`, showing that every constant field carries the trivial vorticity envelope $B = 0$. But Lean now also proves the obstruction the manuscript actually needs: `not_integrable_kineticEnergyDensity_constantVelocityField`, `not_boundedKineticEnergy_constantVelocityField`, and `not_boundedKineticEnergyUpTo_constantVelocityField` show that on whole-space $\mathbb{R}^3$, every nonzero constant velocity field fails the bounded-energy slot. So the tempting “nonzero constant witness” ansatz is ruled out by a real proved theorem rather than by informal handwaving. Lean now also proves the stronger initial-data obstruction the manuscript actually needs for the explicit theorem surfaces: if a field merely matches nonzero constant initial data at time $t = 0$, then its initial kinetic-energy density already agrees with the forbidden constant one. Concretely, `not_integrable_kineticEnergyDensity_zero_of_matchesInitialVelocity_constantInitialVelocity`, `not_boundedKineticEnergyUpTo_of_matchesInitialVelocity_constantInitialVelocity`, and `not_boundedKineticEnergy_of_matchesInitialVelocity_constantInitialVelocity` push the obstruction from the special ansatz $u(t, x) \equiv c$ to any field with nonzero constant initial slice. Consequently `not_nonempty_ExplicitFiniteTimeRegularityWitness_constantInitialVelocity`, `not_ExplicitConcreteNavierStokesRegularityClause_constantInitialVelocity` show that the current explicit witness, global-output, and regularity-clause surfaces all honestly reject nonzero constant initial data on $\mathbb{R}^3$. Lean now also contains the first genuinely nonconstant version of the same obstruction. The vector-calculus layer introduces `linearShearMap`, `linearShearInitialVelocity`, and the actual space-time field `linearShearVelocityField`, proving that the profile $u(t, x) = (ax_1, 0, 0)$ is smooth, divergence-free, has zero time derivative, zero convection, zero Laplacian, and constant vorticity $(0, 0, -a)$. So this is not merely a tempting initial datum: it is an honest

exact Navier–Stokes differential-side candidate on $\mathbb{R} \times \mathbb{R}^3$ with zero pressure. The uniform-vorticity layer then proves `uniformVorticityBoundUpTo_linearShearVelocityField`, so this nonzero candidate carries the explicit bound $B = |a|$. Lean now packages the whole point directly as `linearShearVelocityField_exhibits_concreteCandidate_except_boundedEnergy`: already on the global explicit theorem surface, the steady shear field gives an honest smooth incompressible zero-pressure Navier–Stokes candidate whose only missing slot is bounded kinetic energy. Lean then packages the resulting global-surface mismatch as `linearShearInitialVelocity_exhibits_concreteCandidate_with` the explicit candidate exists, but `ExplicitConcreteNavierStokesGlobalOutput` is still false. Lean now pushes this all the way to the theorem surface: `not_ExplicitConcreteNavierStokesMillenniumTarget` and `not_ConcreteNavierStokesMillenniumTarget` show that the current formal target is outright false, because it quantifies over smooth divergence-free data like nonzero linear shear without imposing a finite-energy admissibility condition on the input side. Lean now also records the repaired whole-space theorem surface itself: *This is a bug in our local over-broad whole-space surrogate, not in the upstream Clay-style theorem surface, and not by itself in Goertzel’s current manuscript route.* The reference Lean repo already builds the $\mathbb{R}^3$ statement with bounded-energy admissibility on the input side, and Goertzel’s draft works with narrower input surfaces such as the torus and Schwartz-type whole-space data. Accordingly, the negative theorem here should be read as a regression guard on our theorem shaping: any local $\mathbb{R}^3$ target that forgets that input restriction is wrong. `finiteInitialKineticEnergy`, `FiniteEnergyAdmissibleNavierStokesProblemData`, `FiniteEnergyConcreteNavierStokesMillenniumTarget`, `ExplicitFiniteEnergyAdmissibleNavierStokesReg` and `FiniteEnergyConcreteNavierStokesMillenniumTarget_iff_explicit`. So the correction is no longer only informal commentary: it is a separate Lean target. Lean now also records the direct theorem-surface implication layer: `ExplicitConcreteNavierStokesRegularityClause_implies_finiteEnergy`, `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_iff_of_finiteInitialKineticEnergy`, `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergy`, and `ConcreteNavierStokesMille` so the repaired theorem surface is now explicitly packaged as a restriction of the older unrepaired one. The zero datum satisfies the repaired input hypothesis via `finiteInitialKineticEnergy_zero`, and `NavierStokesSchwartzData.lean` now adds the manuscript-compatible restricted whole-space input lane: `NSSchwartzInitialVelocity`, `smoothInitialVelocityData_of_schwartzInitialVelocity`, `finiteInitialKineticEnergy_of_schwartzInitialVelocity`, `ExplicitSchwartzAdmissibleNavierStokesReg` and `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_schwartzTarget` package the fact that Schwartz initial data automatically satisfy the repaired $\mathbb{R}^3$ input hypotheses. This gives the local Lean lane a theorem surface that matches the manuscript’s intended torus-to-Schwartz exhaustion story more closely than the old unrepaired whole-space surrogate. `NavierStokesEquationTarget.lean` was also repaired at this stage so the underlying smoothness predicates are genuinely C^∞ rather than the stronger ω -analytic `ContDiff` $\top$ tier; this matters for the Schwartz route because generic Schwartz data are smooth, but not automatically analytic. The next refinement is now present explicitly in `NavierStokesSchwartzDivergenceFreeData.lean`: Lean packages the manuscript’s Appendix H.1 input class $S_\sigma(\mathbb{R}^3)$ itself as the subtype `NSSchwartzDivergenceFreeInitial` of Schwartz data whose initial divergence vanishes, and theorems such as `ExplicitSchwartzDivergenceFreeNavier` and `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_schwartzDivergenceFreeTa` show that the repaired whole-space target now specializes directly to that manuscript-style input surface rather than only to the looser “Schwartz plus a separate divergence-free hypothesis” packaging. The same file now also packages the structured side of that restriction: a manuscript-style datum with $\nu > 0$ canonically produces repaired `FiniteEnergyAdmissibleNavierStokesProblemData`, and `StructuredSchwartzDivergenceFreeNavierStokesRegularityClause_iff` identifies the resulting record-based whole-space clause with the explicit $S_\sigma(\mathbb{R}^3)$ clause. So the manuscript-compatible subtype is now connected to both local theorem surfaces, not only the explicit one.

`NavierStokesSchwartzPartialPeriodization.lean` now adds the next bottom-up Appendix H.2/H.3 precursor: a finite lattice-sum periodization layer. Instead of jumping straight to the full infinite torus-periodization series, Lean first packages the local algebra that the later exhaustion argument needs: `periodizationShift`, `translateInitialVelocity`, and `finitePartialPeriodizedInitialVelocity`, together with proofs that spatial translation preserves smoothness and initial divergence, finite sums of translated data preserve both properties, and finite partial periodizations therefore preserve the manuscript's $S_\sigma(\mathbb{R}^3)$ input class via `finitePartialPeriodizationSchwartzDivergenceFreeInitialVelocity`. The companion regression file checks the empty and singleton-zero partial periodizations explicitly, so the torus-to-whole-space route now has a verified finite prelude rather than only an informal next-step note. `NavierStokesSchwartzPeriodizationBoxes.lean` now adds the next H.4-side exhaustion layer on top of that prelude: the concrete centered cubic boxes $[-N, N]^3 \subset \mathbb{Z}^3$, the shell decomposition from radius N to radius $N + 1$, and the boxed partial-periodization operators obtained by specializing the general finite-sum interface to those boxes. This makes the truncation family itself available as named Lean data via `centeredLatticeBox`, `centeredLatticeShell`, and `boxedPartialPeriodizedInitialVelocity`, with direct zero-radius and successor-step theorems. So the manuscript's large-box exhaustion is no longer just an intended indexing convention: the finite approximation family and its one-step growth law now exist explicitly in the local formalization. `NavierStokesSchwartzPeriodizationConvergence.lean` adds the first honest local limit interface above those finite boxes. Instead of claiming the full Schwartz periodization limit already exists, Lean now packages cube-local convergence of the boxed partial periodizations on $Q_R = [-R, R]^3$ via `BoxedPartialPeriodizationConvergesOnCube`, proves monotonicity in the cube radius, and derives the key structural consequence that if the boxed sums converge on a fixed cube then the added shell contributions must converge to zero there. The newest strengthening also packages this shell-vanishing corollary directly for genuine Schwartz data and nonzero lattice period, so Appendix H.4's first local decay step is now available in Lean without first restating a separate convergence hypothesis. This is exactly the first nontrivial convergence fact needed before any Arzelà–Ascoli / compactness step in Appendix H.4 can be stated cleanly. `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_zero` shows the repaired explicit clause is inhabited there. On the other hand, Lean now also records that this repaired explicit clause is still vacuous outside the finite-energy domain: `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_generic` packages the generic logical fact, while `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_constantInitialVelocity` and `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_linearShearInitialVelocity_withConcreteData` show concretely that the repaired clause can still be true on the old bad families while the unrepaired concrete regularity clause is false there. `not_finiteInitialKineticEnergy_linearShearInitialVelocity` and `not_finiteInitialKineticEnergy_constantInitialVelocity` show that the counterexample families are now excluded at the input side rather than only killed later by the output bounded-energy slot, and Lean now also records that this exclusion is literal on the repaired structured theorem surface via the generic theorem `not_nonempty_FiniteEnergyAdmissibleNavierStokesProblemData_of_not_finiteInitialKineticEnergy` with the concrete constant and shear instances exposed as `not_nonempty_FiniteEnergyAdmissibleNavierStokesProblemData_constantInitialVelocity` and `not_nonempty_FiniteEnergyAdmissibleNavierStokesProblemData_linearShearInitialVelocity`. Lean now also packages the positive converse for fixed concrete data: `nonempty_FiniteEnergyAdmissibleNavierStokesProblemData_of_finiteInitialKineticEnergy` says that once viscosity positivity, initial smoothness, and initial divergence-freeness are fixed, inhabiting the repaired structured input surface is equivalent to having finite initial kinetic energy. At the clause level, `finiteEnergyConcreteNavierStokesGlobalRegularityClause_iff` now makes the same repaired surface explicit: for a fixed finite-energy datum, the repaired structured fully concrete clause is just the repaired explicit regularity clause packaged through the concrete target shell. Lean now also packages the repaired/unrepaired structured-clause equivalence on that same fixed datum directly as `finiteEnergyConcreteNavierStokesGlobalRegularityClause_iff_of_finiteInitialKineticEnergy`.

so the repaired concrete shell is no longer only compared to the unrepaired one through the explicit surface. The equation layer now also exposes the forgetful directions themselves: `ExplicitFiniteEnergyAdmissibleNavierStokesGlobalRegularityClause` packages the repaired explicit-to-unrepaired explicit step on a fixed finite-energy datum, while `finiteEnergyConcreteNavierStokesGlobalRegularityClause_implies_concreteNavierStokesGlobalRegularityClause` and `concreteNavierStokesGlobalRegularityClause_implies_finiteEnergyConcreteNavierStokesGlobalRegularityClause` do the same directly on the structured concrete shell. At the theorem surface, Lean now also packages the structured-to-explicit export directly as `ConcreteNavierStokesMillenniumTarget_implies_ExplicitConcreteNavierStokesMillenniumTarget` and `FiniteEnergyConcreteNavierStokesMillenniumTarget_implies_ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget` so downstream continuation files can reuse these bridges instead of reopening the corresponding `_iff_explicit` equivalences locally. It now also packages the mixed bridge `ConcreteNavierStokesMillenniumTarget_implies_finiteEnergyConcreteNavierStokesMillenniumTarget` so downstream repaired continuation exports no longer need to reconstruct the two-step path through the repaired structured theorem surface either. The continuation layer now also uses that repaired input condition analytically: `ExplicitFiniteTimeRegularityWitness.finiteInitialKineticEnergy` and `not_nonempty_ExplicitFiniteTimeRegularityWitness_of_not_finiteInitialKineticEnergy` show that any actual slab witness on a nonnegative horizon already forces finite initial energy. Lean now packages that fact directly into repaired continuation surfaces: `ExplicitFiniteEnergyUniformVorticityContinuationTarget`, `ExplicitFiniteEnergyUniformVorticityContinuationClause`, `ExplicitFiniteEnergyBKMContinuationTarget`, and `ExplicitFiniteEnergyBKMContinuationClause`. The equivalence theorems `ExplicitFiniteEnergyUniformVorticityContinuationTarget_implies_ExplicitFiniteEnergyBKMContinuationTarget` and `ExplicitFiniteEnergyBKMContinuationClause_implies_of_nonneg_horizon` show that on non-negative slabs these repaired clauses coincide with the older ones, precisely because any real slab witness already carries finite initial energy. Consequently the repaired theorem target now lands directly in both repaired continuation targets via `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_ExplicitFiniteEnergyUniformVorticityContinuationTarget` and `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_finiteEnergyBKMContinuationTarget` while the older unrepaired whole-space theorem surfaces now also reach these same repaired continuation targets directly through `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationTarget`, `ConcreteNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationTarget`, `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergyBKMContinuationTarget`, and `ConcreteNavierStokesMillenniumTarget_implies_finiteEnergyBKMContinuationTarget`, and the same fixed-horizon directness is now exposed at clause level via `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationClause`, `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationClause`, `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_finiteEnergyBKMContinuationClause`, and `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergyBKMContinuationClause`, while the older unrepaired continuation layer now also exposes the analogous target-level admissibility lifts `ExplicitUniformVorticityContinuationTarget_implies_finiteEnergyUniformVorticityContinuationTarget` and `ExplicitBKMContinuationTarget_implies_finiteEnergyBKMContinuationTarget`, as well as the fixed-horizon corollaries `ExplicitUniformVorticityContinuationClause_implies_finiteEnergy` and `ExplicitBKMContinuationClause_implies_finiteEnergy`, so the same unrepaired→repaired step is now exposed uniformly on both the target and clause surfaces. The unrepaired continuation layer also exposes the analogous fixed-horizon corollaries `ExplicitConcreteNavierStokesMillenniumTarget_implies_uniformVorticityContinuationClause`, `ConcreteNavierStokesMillenniumTarget_implies_uniformVorticityContinuationClause`, `ExplicitConcreteNavierStokesMillenniumTarget_implies_BKMContinuationClause`, and `ConcreteNavierStokesMillenniumTarget_implies_BKMContinuationClause`. The uniform side now also exposes the lower-level clause repackaging bridges `ExplicitConcreteNavierStokesRegularityClause_implies_ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause` and `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_implies_ExplicitConcreteNavierStokesRegularityClause` as well as their fixed-datum all-horizons families `ExplicitConcreteNavierStokesRegularityClause_implies_ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause` and `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_implies_ExplicitConcreteNavierStokesRegularityClause` so both fixed-horizon uniform clauses and their fixed-datum target families can be obtained directly from the imported whole-space explicit clause surfaces. The theorem surface now re-exports those fixed-datum families as `ExplicitConcreteNavierStokesMillenniumTarget_implies_uniformVorticityContinuationClause`.

and `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationClause_allHorizons` so downstream code can stay on the theorem-target API while still recovering the full $\forall T$ continuation-clause family without reopening the whole-space clause binder locally. The same pattern now also covers the structured and repaired theorem surfaces via `FiniteEnergyConcreteNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationClause_allHorizons`, `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationClause_allHorizons`, and `ConcreteNavierStokesMillenniumTarget_implies_uniformVorticityContinuationClause_allHorizons`. The BKM side now also exposes the lower-level clause repackaging bridges `ExplicitConcreteNavierStokesRegularClause_implies_finiteEnergyBKMContinuationClause_allHorizons` and `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_implies_ExplicitFiniteEnergyBKMContinuationClause_allHorizons` as well as their fixed-datum all-horizons families `ExplicitConcreteNavierStokesRegularityClause_implies_ExplicitFiniteEnergyBKMContinuationClause_allNonnegHorizons` and `ExplicitFiniteEnergyAdmissibleNavierStokesRegularityClause_implies_ExplicitFiniteEnergyBKMContinuationClause_allNonnegHorizons` so both fixed-horizon BKM clauses and their fixed-datum target families can be obtained directly from the imported whole-space explicit clause surfaces. The theorem surface now re-exports those fixed-datum families as `ExplicitConcreteNavierStokesMillenniumTarget_implies_BKMContinuationClause_allNonnegHorizons` and `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_finiteEnergyBKMContinuationClause_allNonnegHorizons` so the same direct $\forall T$ access is available on the explicit BKM theorem-target API as well. The same pattern now also covers the structured and repaired theorem surfaces via `FiniteEnergyConcreteNavierStokesMillenniumTarget_implies_finiteEnergyBKMContinuationClause_allNonnegHorizons`, `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergyBKMContinuationClause_allNonnegHorizons`, `ConcreteNavierStokesMillenniumTarget_implies_BKMContinuationClause_allNonnegHorizons`, and the zero-data BKM clause baselines now factor through those reusable bridges instead of reopening the witness and envelope binders locally. The repaired continuation layer itself now contains the same BKM-to-uniform bridge as the unrepaired one: `ExplicitFiniteEnergyBKMContinuationClause_implies_finiteEnergyUniformVorticityContinuationClause_allNonnegHorizons` and `ExplicitFiniteEnergyBKMContinuationTarget_implies_finiteEnergyUniformVorticityContinuationClause_allNonnegHorizons` to transport repaired BKM data directly to repaired uniform-vorticity data, and the BKM target surface now also re-exports the corresponding fixed-datum uniform clause families as `ExplicitBKMContinuationTarget_implies_finiteEnergyUniformVorticityContinuationClause_allNonnegHorizons`, `ExplicitBKMContinuationTarget_implies_finiteEnergyUniformVorticityContinuationClause_allNonnegHorizons`, and `ExplicitFiniteEnergyBKMContinuationTarget_implies_finiteEnergyUniformVorticityContinuationClause_allNonnegHorizons` so downstream code can recover the entire unrepaired or repaired uniform clause family directly from a BKM target assumption without first passing through the target-valued bridge theorem, while the same file now also records the composed theorem-surface exports `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationClause_allNonnegHorizons`, `ExplicitConcreteNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationClause_allNonnegHorizons`, their repaired/concrete analogues, and the matching fixed-datum clause families such as `ExplicitConcreteNavierStokesRegularityClause_implies_finiteEnergyBKMContinuationClause_allNonnegHorizons` and `ConcreteNavierStokesMillenniumTarget_implies_finiteEnergyUniformVorticityContinuationClause_allNonnegHorizons` so the BKM route remains explicit even at the whole-theorem and whole-family surface, while the older bridge theorems to the unrepaired clauses survive as corollaries on nonnegative horizons. These older nonnegative-horizon forgetful steps are now also packaged as direct theorem-target families `ExplicitFiniteEnergyUniformVorticityContinuationTarget_implies_uniformVorticityContinuationClause_allNonnegHorizons`, `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_uniformVorticityContinuationClause_allNonnegHorizons`, `FiniteEnergyConcreteNavierStokesMillenniumTarget_implies_uniformVorticityContinuationClause_allNonnegHorizons`, `ExplicitFiniteEnergyBKMContinuationTarget_implies_BKMContinuationClause_allNonnegHorizons`, `ExplicitFiniteEnergyAdmissibleNavierStokesMillenniumTarget_implies_BKMContinuationClause_allNonnegHorizons`, `FiniteEnergyConcreteNavierStokesMillenniumTarget_implies_BKMContinuationClause_allNonnegHorizons` and `ExplicitFiniteEnergyBKMContinuationTarget_implies_uniformVorticityContinuationClause_allNonnegHorizons` so downstream code can use a single repaired target and then recover the entire “for every nonnegative slab” family without reopening the fixed-horizon forgetful lemma by hand. These repaired targets are also already operational on the concrete zero family: `ExplicitFiniteEnergyUniformVorticityContinuationTarget_implies_zeroGlobalOutput` and `ExplicitFiniteEnergyBKMContinuationTarget_implies_zeroGlobalOutput` apply them di-

rectly to the canonical zero slab witness together with the now explicit input fact `finiteInitialKineticEnergy_zero`. The API is now also clean one layer lower: `ExplicitFiniteEnergyUniformVorticityContinuationClause_implies_zeroGlobalOutput` and `ExplicitFiniteEnergyBKMContinuationClause_implies_zeroGlobalOutput` show that the repaired clause surface itself already has the same concrete zero-data operational consequence as the unrepaired one. The same is now true for the time-gauge sanity check: `ExplicitUniformVorticityContinuationTarget_implies_zeroGlobalOutput_timeOnlyPressure` and `ExplicitBKMContinuationTarget_implies_zeroGlobalOutput_timeOnlyPressure` now package the unrepaired target-level version directly, while `ExplicitFiniteEnergyUniformVorticityContinuationClause_implies_zeroGlobalOutput_timeOnlyPressure` and `ExplicitFiniteEnergyBKMContinuationClause_implies_zeroGlobalOutput_timeOnlyPressure` show that the repaired clause layer also survives adding any smooth time-dependent pressure gauge on a larger zero slab and restricting back down. Lean now also packages the constant-pressure special cases directly on all four continuation surfaces via `ExplicitUniformVorticityContinuationClause_implies_zeroGlobalOutput_constantPressure`, `ExplicitUniformVorticityContinuationTarget_implies_zeroGlobalOutput_constantPressure`, `ExplicitBKMContinuationClause_implies_zeroGlobalOutput_constantPressure`, and `ExplicitBKMContinuationTarget_implies_zeroGlobalOutput_constantPressure`, plus their repaired finite-energy analogues. The same zero-data operational API now exists one layer higher on the actual whole-space theorem surfaces: `ExplicitConcreteNavierStokesMillenniumTarget_implies_zeroGlobalOutput`, `ConcreteNavierStokesMillenniumTarget_implies_zeroGlobalOutput`, and their repaired finite-energy analogues. That whole-space layer is now also normalized through the same gauge interface: `ExplicitConcreteNavierStokesMillenniumTarget_implies_zeroGlobalOutput_addPressureOfZeroSpatialGradient`, `ConcreteNavierStokesMillenniumTarget_implies_zeroGlobalOutput_addPressureOfZeroSpatialGradient`, and the repaired finite-energy analogues package the generic zero-spatial-gradient transport once, and the older time-only and constant-pressure theorem-surface exports are now just wrappers around that shared zero-data theorem. The repaired target surface now packages the same fact directly via `ExplicitFiniteEnergyUniformVorticityContinuationTarget_implies_zeroGlobalOutput_timeOnlyPressure` and `ExplicitFiniteEnergyBKMContinuationTarget_implies_zeroGlobalOutput_timeOnlyPressure`. The continuation layer then packages the whole point directly as `linearShearVelocityField_exhibits_uniformCauchy`, there exist explicit u, p, B satisfying the smoothness, momentum, incompressibility, initial-condition, and uniform-vorticity parts of the finite-time witness surface, and the only missing slot is bounded kinetic energy. The BKM file then sharpens the same point on the stricter time-integrated surface: `linearShearVelocityField_has_constantBKMEvelope` proves that the same field carries the literal constant envelope $\Omega(t) = |a|$ with integral bound $T|a|$, and `linearShearVelocityField_exhibits_BKMCandidate` packages it as a full BKM-side candidate whose only missing slot is again bounded kinetic energy. Lean now also packages the precise mismatch directly as `linearShearInitialVelocity_exhibits_BKMCandidate`, the explicit BKM-side candidate exists, but the finite-time witness type is still empty. The continuation file also defines the associated explicit density `linearShearKineticEnergyDensity` and proves `not_integrable_linearShearKineticEnergyDensity` by an exact scaling argument on $\mathbb{R}^3$: if the density were integrable, Haar scaling by $x \mapsto 2x$ would force its integral to be both $4I$ and $I/8$, hence zero, contradicting positivity. This is then pushed from the special shear ansatz to any field matching nonzero linear-shear initial data at $t = 0$ via `not_integrable_kineticEnergyDensity_zero_of_matchesInitialVelocity`, `not_boundedKineticEnergyUpTo_of_matchesInitialVelocity_linearShearInitialVelocity`, and `not_boundedKineticEnergy_of_matchesInitialVelocity_linearShearInitialVelocity`. Consequently `not_nonempty_ExplicitFiniteTimeRegularityWitness_linearShearInitialVelocity`, `not_ExplicitConcreteNavierStokesGlobalOutput_linearShearInitialVelocity`, and `not_ExplicitConcreteNavierStokesGlobalOutput_linearShearInitialVelocity` show that the current explicit bounded-energy theorem surfaces also reject this genuinely nonconstant initial-data family. At the same time, Lean now exposes a precise target-shape caveat: the current continuation clauses themselves remain true on such impossible data, but only vacuously, because they quantify over a finite-time witness type that is already empty there. This is now stated explicitly by `ExplicitUniformVorticityContinuationClause_of_not_nonempty_finiteTimeWitness` and `ExplicitBKMContinuationClause_of_not_nonempty_finiteTimeWitness`, which isolate the

vacuity mechanism itself, and then concretely by `ExplicitUniformVorticityContinuationClause_constantInitialVelocity` and `ExplicitBKMContinuationClause_constantInitialVelocity`; Lean now also shows the same mismatch on the nonconstant shear family via `ExplicitUniformVorticityContinuationClause_linearShearInitialVelocity` and `ExplicitBKMContinuationClause_linearShearInitialVelocity`; and it packages both mismatches directly as `ExplicitUniformVorticityContinuationClause_constantInitialVelocity_without_regularity` and `ExplicitBKMContinuationClause_constantInitialVelocity_without_regularity`, as well as `ExplicitUniformVorticityContinuationClause_linearShearInitialVelocity_without_regularity` and `ExplicitBKMContinuationClause_linearShearInitialVelocity_without_regularity`. The repaired continuation clauses inherit a second vacuity mode from the finite-energy input filter itself: `ExplicitFiniteEnergyUniformVorticityContinuationClause_of_not_finiteInitialKineticEnergy` and `ExplicitFiniteEnergyBKMContinuationClause_of_not_finiteInitialKineticEnergy` package the generic logical fact, while the constant and shear families are now also recorded directly on those repaired surfaces by `ExplicitFiniteEnergyUniformVorticityContinuationClause_constantInitialVelocity`, `ExplicitFiniteEnergyBKMContinuationClause_constantInitialVelocity_without_regularity`, `ExplicitFiniteEnergyUniformVorticityContinuationClause_linearShearInitialVelocity_without_regularity`, and `ExplicitFiniteEnergyBKMContinuationClause_linearShearInitialVelocity_without_regularity`. So the present continuation surfaces are not by themselves existence claims; the zero family supplies one honest concrete nonvacuity route, and linear shear now supplies an honest nonzero PDE-side candidate with explicit vorticity control, but genuinely nonzero inhabited finite-time witnesses still fail here precisely at the bounded-energy slot. This again isolates a defect in the present local whole-space continuation surface; it does not by itself rule out reasonable torus or Schwartz-restricted adaptations of the Goertzel route. On the positive API side, the same file now also packages the repaired-to-unrepaired uniform forgetful step on nonnegative slabs directly as `ExplicitFiniteEnergyUniformVorticityContinuationClause_implies_uniformVorticityContinuationClause` and its target-to-fixed-horizon corollary `ExplicitFiniteEnergyUniformVorticityContinuationTarget_implies_uniformVorticityContinuationTarget`, so the repaired whole-space theorem-surface export on a fixed nonnegative horizon can now factor through one reusable continuation-target theorem instead of reopening the repaired/unrepaired equivalence locally. The repaired uniform clause now also inherits the same restriction-based horizon monotonicity API directly: `ExplicitFiniteEnergyUniformVorticityContinuationClause_apply_of_uniformVorticityContinuationClause` and `ExplicitFiniteEnergyUniformVorticityContinuationClause_mono_horizon` show that once finite initial energy is fixed, the repaired continuation surface restricts along shorter slabs exactly as the unrepaired one does.

The next file, `NavierStokesBKMContinuationTarget.lean`, sharpens that surface toward the actual BKM shape. Instead of a uniform-in-time bound, it asks for an explicit scalar envelope $\Omega : \mathbb{R} \rightarrow \mathbb{R}$ such that $\|\omega(t, x)\| \leq \Omega(t)$ on the slab, $\Omega(t) \geq 0$ on $[0, T]$, and Ω is integrable on $[0, T]$ with an explicit integral bound. This is still only a target, not a proof, but it is a more theorem-like continuation surface than the earlier uniform bound. Lean again proves the sanity theorem that the full explicit global theorem target subsumes this BKM-style continuation target, and hence so does the structured fully concrete theorem target. In this file we now also have a nontrivial bridge theorem: a uniform slab vorticity bound gives an explicit constant-in-time envelope that is integrable on $[0, T]$, and therefore the BKM-style continuation target implies the uniform-vorticity continuation target. So this is no longer only parallel one-way subsumption from the full global theorem target; there is now a real continuation-surface implication that uses the continuation hypotheses. Combining that implication with the new zero-data uniform instantiation gives the corresponding operational BKM corollary: under the BKM target assumption and positive viscosity, one gets an explicit global output for zero initial data. At the same layer, Lean now also contains the direct clause-level operational formulation `ExplicitBKMContinuationClause_implies_zeroGlobalOutput`, which applies one concrete BKM clause instance directly to the canonical zero slab witness and explicit envelope data. The enve-

lope input is now itself proved directly in `zeroFiniteTimeRegularityWitness_has_BKMEnvelope`, with the explicit choice $\Omega(t) = 0$ and $B = 0$ on $[0, T]$, via reusable zero-envelope components `vorticityEnvelopeOn_zero` and `integrableVorticityEnvelopeOn_zero`. It now also contains the constant-envelope lemma `integrableVorticityEnvelopeOn_const`, which on nonnegative slabs $[0, T]$ turns the constant envelope value B into the explicit linear bound $T \cdot B$, and the corresponding bridge `uniformVorticityBoundUpTo_implies_constantBKMEnvelope` showing that a uniform slab vorticity bound already produces that exact finite-time constant-envelope package. The same file now also adds reusable time-restriction transport lemmas `vorticityEnvelopeOn_restrict` and `integrableVorticityEnvelopeOn_restrict`, and uses them to refactor the uniform-to-BKM restriction bridge `uniformVorticityBoundUpTo_implies_BKMEnvelope_restrict`. At the same layer it now also includes a clause-level bridge theorem `ExplicitBKMContinuationClause_implies_uniformVorticity` and the witness-level restricted-envelope transport theorem `ExplicitFiniteTimeRegularityWitness.BKMEnvelope`. Lean now also packages the repaired-to-unrepaired BKM forgetful bridge on nonnegative slabs: `ExplicitFiniteEnergyBKMContinuationClause_implies_BKMContinuationClause_of_nonneg_horizon` and its target-to-fixed-horizon corollary `ExplicitFiniteEnergyBKMContinuationTarget_implies_BKMContinuation`. It then composes that forgetful step with the unrepaired BKM-to-uniform bridge to obtain the repaired-to-unrepaired uniform variant on the same slabs: `ExplicitFiniteEnergyBKMContinuationClause_implies_uniformVorticity` and its target-to-fixed-horizon corollary `ExplicitFiniteEnergyBKMContinuationTarget_implies_uniformVorticity`. The repaired BKM clause now also has the direct restriction-based horizon API `ExplicitFiniteEnergyBKMContinuationClause_restrict` and `ExplicitFiniteEnergyBKMContinuationClause_mono_horizon`, so the repaired BKM surface itself now transports to larger horizons by the same witness-restriction argument as the unrepaired clause. The bridge stack now also reaches the repaired uniform surface directly: `ExplicitBKMContinuationClause_implies_finiteEnergyUniformVorticityContinuationClause` records the same-horizon unrepaired-BKM to repaired-uniform step, while `ExplicitBKMContinuationClause_restrict` and `ExplicitFiniteEnergyBKMContinuationClause_implies_finiteEnergyUniformVorticityContinuationClause` package the corresponding shorter-slab to larger-slab transport there. At the theorem-surface level this now yields the direct bridge `ExplicitBKMContinuationTarget_implies_finiteEnergyUniformVorticityContinuationTarget` in addition to the older unrepaired-uniform and repaired-uniform target exports. The repaired theorem-surface fixed-horizon BKM export now factors through that target-level forgetful theorem, rather than reopening the repaired/unrepaired BKM equivalence locally. The newest pass sharpens that witness-level transport by using envelope nonnegativity on the restricted slab, via `vorticityEnvelopeOn.integrable_restrict`, so the witness-level bound is the restricted integral itself rather than its absolute value. The same pass also adds the BKM clause horizon-monotonicity theorem `ExplicitBKMContinuationClause_mono_horizon`. The latest bottom-up addition to that file is a small additive algebra for concrete envelopes: `vorticityEnvelopeOn_add` combines two slab envelopes into an envelope for the summed field $u + v$, using the concrete curl linearity lemma `spatialVorticity_add`, and `integrableVorticityEnvelopeOn_add` adds the corresponding time-integral bounds. So the BKM layer now supports not just restriction and constantization of envelope data, but also explicit composition of envelope data under field addition. The newest pass makes that addition theorem witness-friendly: `vorticityEnvelopeOn_add_of_smooth` discharges the pointwise differentiability side condition from full space-time smoothness, and the same file now also contains the direct sum-bridge corollaries `uniformVorticityBoundUpTo_add_implies_BKMEnvelope` and `uniformVorticityBoundUpTo_add_implies_constantBKMEnvelope`, which turn two smooth uniform slab bounds into explicit BKM-envelope data for the summed field. The next pass adds the rescaling/sign-flip side of the same bottom-up algebra: `vorticityEnvelopeOn_const_smul` transports a slab envelope along constant velocity rescaling $u \mapsto a \cdot u$ with the expected factor $|a| \Omega(t)$, while `integrableVorticityEnvelopeOn_const_mul` and `integrableVorticityEnvelopeOn_abs_const_smul` transport the time-integrability bound itself under nonnegative and absolute-value rescaling. The

sign-flip corollary `vorticityEnvelopeOn_neg` then comes for free. So the concrete BKM layer now supports not only restriction, constantization, and addition, but also explicit scaling symmetry of envelope data. The newest pass then turns that algebra into subtraction theorems: `vorticityEnvelopeOn_sub_of_smooth` gives a concrete envelope for $u - v$ from smooth envelopes for u and v , and the bridge corollaries `uniformVorticityBoundUpTo_sub_implies_BKMEnvelope` and `uniformVorticityBoundUpTo_sub_implies_constantBKMEnvelope` package the same perturbative move at the uniform-to-BKM interface. So the bottom-up continuation algebra now handles both addition and subtraction of smooth slab data. The latest pass then bundles the envelope and integrability sides back together into reusable package theorems `BKMEnvelopeData_add_of_smooth` and `BKMEnvelopeData_sub_of_smooth`. These no longer just say that the pointwise bound or the time-integral bound can be transported separately; they package exact finite-slab BKM data for $u + v$ and $u - v$ in one theorem surface, which is a better mouth for later witness-level perturbative arguments. The newest pass then lifts those bundled theorems to the witness layer itself via `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_add` and `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_sub` if two finite slab witnesses each carry explicit BKM data on the same horizon, then the sum or difference of their velocity fields carries explicit BKM data too. That is still not a continuation proof, but it is a cleaner perturbative surface for future witness-level arguments than re-unpacking the smoothness and envelope data by hand each time. The newest pass adds the rescaling/sign-flip side at the same witness level: `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_const_smul` transports explicit BKM data along constant velocity rescaling, and `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_neg` specializes that to sign reversal. So the witness-level perturbative surface now supports addition, subtraction, scaling, and negation of the carried BKM packages, without claiming that the transformed velocity fields are themselves full Navier–Stokes witnesses. The newest pass then combines perturbation with horizon restriction directly: `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_add_restrict` and `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_sub_restrict` show that once the sum or difference field carries explicit BKM data on a slab $0 \leq t \leq T$, the same concrete package can be pushed down to any shorter slab $0 \leq t \leq T' \leq T$ without rebuilding the envelope algebra by hand. The newest pass closes the same smaller-slab transport under scaling symmetry as well: `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_const_smul_restrict` and `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_neg_restrict` show that constant rescaling and sign flip commute with this witness-level restriction surface too. The newest pass packages the whole perturbative pattern into one direct API theorem: `BKMEnvelopeData_linearCombination_of_smooth` gives exact finite-slab BKM data for $a \bullet u + b \bullet v$, with the envelope $|a| \Omega + |b| \Omega'$ and the corresponding bound $|a| B + |b| B'$, and `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_linearCombination_restrict` pushes that same linear-combination package down to any shorter slab $0 \leq t \leq T' \leq T$. So later perturbative arguments can now consume one concrete linear-combination theorem instead of manually chaining scale, add, and restrict lemmas. The newest pass mirrors this on the uniform-vorticity side too: `uniformVorticityBoundUpTo_linearCombination` packages the same $|a| B + |b| B'$ bound for the concrete uniform criterion, and the bridge corollaries `uniformVorticityBoundUpTo_linearCombination_implies_BKMEnvelope` and `uniformVorticityBoundUpTo_linearCombination_restrict` push that uniform perturbative surface directly back into the BKM layer. The newest pass lifts this to the witness layer on both sides: `ExplicitFiniteTimeRegularityWitness.uniformVorticityBound_linearCombination` and `ExplicitFiniteTimeRegularityWitness.uniformVorticityBound_linearCombination_restrict` give direct witness-level uniform control for linear combinations on the full slab and on shorter slabs, while `uniformVorticityBoundUpTo_linearCombination_implies_BKMEnvelope_restrict` now gives the corresponding field-level restricted BKM bridge directly, so the witness theorem can sit on a concrete PDE lemma instead of restitching the conversion by hand. Then `ExplicitFiniteTimeRegularityWitness.uniformVorticityBound_linearCombination_implies_BKMEnvelope`

turns that same smaller-slab witness-level uniform control directly into explicit BKM data. The newest pass then adds the constant-envelope version of that bridge on nonnegative shorter slabs: `uniformVorticityBoundUpTo_implies_constantBKMEnvelope_restrict` packages the restricted constant envelope $\Omega(t) = B$, and `uniformVorticityBoundUpTo_linearCombination_implies_constantBKMEnvelope_restrict` does the same for the linear-combination surface itself, while `ExplicitFiniteTimeRegularityWitness.uniformVorticityBoundUpTo_restrict` lifts the same constant-envelope route to witness-level linear combinations. The generic BKM side now has the matching restricted package too: `BKMEnvelopeData_linearCombination_restrict_of_smooth` gives the shorter-slab linear-combination envelope data directly at the concrete PDE layer, and `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_linearCombination_restrict` is now just the witness-level lift of that same bottom-up package. The continuation layer now consumes these restriction bridges honestly on actual witnesses: `ExplicitUniformVorticityContinuationClause_apply_of_uniformVorticityBoundUpTo_restrict` applies a smaller-slab uniform clause to a larger-slab witness after restriction, `ExplicitBKMContinuationClause_apply_of_restrict` does the same for explicit BKM data, and `ExplicitBKMContinuationClause_apply_of_uniformVorticityBoundUpTo_restrict` feeds the restricted uniform-to-BKM bridge directly into the BKM clause. This keeps the continuation side honest: it uses smaller-slab restrictions of genuine witnesses, rather than treating linear combinations of unrelated witnesses as new Navier–Stokes solutions. The latest pass also makes this continuation surface explicitly pressure-gauge invariant: `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_addTimeOnlyPressure_restrict` shows that the carried BKM package is unchanged after adding any smooth time-only pressure gauge to the witness, and `ExplicitBKMContinuationClause_apply_of_BKMEnvelope_addTimeOnlyPressure_restrict` shows that the BKM clause can therefore be applied after such a gauge change without any hidden pressure normalization assumption. The same invariance now also survives smaller-slab restriction: `ExplicitFiniteTimeRegularityWitness.BKMEnvelope_addTimeOnlyPressure_restrict` pushes the unchanged BKM package down to any shorter slab after the gauge change, and `ExplicitBKMContinuationClause_apply_of_restrict` turns that directly into a smaller-horizon operational clause theorem. Lean now also exposes the same operational mouth starting only from a uniform vorticity bound: `ExplicitBKMContinuationClause_apply_of_uniformVorticityBound_restrict` applies the BKM clause after a same-horizon smooth time-only pressure-gauge change using the unchanged uniform bound, while `ExplicitBKMContinuationClause_apply_of_uniformVorticityBound_addTimeOnlyPressure_restrict` does the same after restricting a larger-slab gauged witness to a shorter horizon. Lean now also instantiates this on the concrete zero witness family: `zeroFiniteTimeRegularityWitness_addTimeOnlyPressure_restrict` shows that the gauged-and-restricted zero witness itself already carries the explicit shorter-slab zero BKM envelope, and `ExplicitBKMContinuationClause_implies_zeroGlobalOutput_timeOnlyPressure_restrict` shows that a smaller-slab BKM clause can therefore be applied directly to that concrete gauged zero witness.

2.23 Compatibility Frontier

In `FeffermanCompatibilityFrontier.lean`, Lean peels that top-down bridge one step further. It defines `AlmostFeffermanBridge`, which contains:

1. a candidate velocity/pressure pair;
2. the corresponding regularity certificate;
3. the corresponding dynamics certificate;
4. the corresponding energy certificate;

but deliberately omits vorticity-to-velocity compatibility.

Lean then proves that once a compatibility proof is supplied, the full `TopDownFeffermanBridge` and hence the local Fefferman-style clause follow immediately. It also adds package-level reduction

theorems showing that `ColeHopfIdentityData`, `FiniteModeColeHopfData`, `ColeHopfKernelSemigroupData`, `FiniteModeColeHopfKernelData`, `ManuscriptTruncationPackage`, and `SharedApproximationPackage` all reduce to that same compatibility frontier.

So the live constructive NS bottleneck is now even sharper in Lean:

once candidate fields and their certificates are fixed, the remaining theorem is exactly the compatibility of the current vorticity tendril with the chosen velocity field.

2.24 Weak Self-Model Realization

In `FeffermanWeakSelfModel.lean`, Lean now proves the top-down stack is not vacuous. The file defines a weak but nontrivial predicate kit where:

1. smoothness, momentum, incompressibility, and initial-condition predicates are left trivial;
2. the only nontrivial output-side predicate is a uniform pointwise bound on the velocity field;
3. compatibility means the chosen velocity field is literally the current vorticity tendril.

With that choice, Lean constructs a canonical self-model:

1. velocity is the tendril vorticity itself;
2. pressure is identically zero;
3. the tendril envelope supplies the boundedness certificate;
4. compatibility is tautological.

It then proves:

1. every `UniformVorticityTendril` realizes the weak local Fefferman-style clause;
2. the same realization can be seen directly through the compatibility frontier;
3. all current NS packages therefore already realize that weak clause.

So the constructive picture is now stronger:

the top-down NS bridge is not merely a decomposition of obligations; for a weak but genuine target, the current grassroots tendril already closes the whole stack in Lean.

2.25 Seeded and Bounded Seeded Strengthenings

In `FeffermanSeededSelfModel.lean`, Lean strengthens the weak self-model in two concrete ways:

1. the pressure predicate is no longer trivial: it must be identically zero;
2. the initial-condition predicate is no longer trivial: velocity at the distinguished base time 1 must match the seed slice extracted from the tendril itself.

The file proves that every current NS package still realizes that stronger seeded clause by the same self candidate and the same compatibility proof.

In `FeffermanBoundedSeededSelfModel.lean`, Lean strengthens the target once more by making the velocity-side regularity predicate nontrivial too: it must satisfy the same uniform pointwise bound already carried by the tendril envelope. The same self candidate still closes the clause.

In `FeffermanPressureSeededSelfModel.lean`, Lean strengthens the target again by making one more dynamics field nontrivial: the momentum slot now also requires the pressure field to vanish. Since the canonical self candidate still has identically zero pressure, the same self candidate and the same compatibility proof still close the clause.

So the current top-down picture is sharper:

the current NS tendril now closes not just a boundedness-only toy target, but a proof-carrying family of self-model targets with zero pressure, explicit initial slice, a nontrivial velocity-bound predicate, and now also a nontrivial momentum-slot pressure condition.

2.26 First Non-Self Compatibility Frontier

In `FeffermanScaledPressureSeededFrontier.lean`, Lean introduces the first genuinely non-self candidate: velocity is a scalar multiple of the current vorticity tendril, while pressure stays identically zero. Lean proves:

1. the same strengthened pressure-seeded target still has full regularity, dynamics, and energy certificates for this rescaled candidate;
2. therefore the same-route version now reaches an explicit `AlmostFeffermanBridge` with a non-trivial compatibility mouth;
3. a nearby descendant route with rescaled compatibility closes immediately;
4. the same-route mouth already closes in special regimes, including the trivial scaling $a = 1$ and the degenerate zero-vorticity case.

So the constructive picture is sharper again:

Lean now contains not only self-model closures, but also the first explicit non-self candidate whose certificates are complete and whose remaining obstacle is exactly a compatibility theorem.

In `FeffermanSeedBlendPressureFrontier.lean`, Lean pushes this one step further. The new non-self candidate has velocity

$$u(t, x) = a \omega(t, x) + b \omega(1, x),$$

mixing the live tendril with its seeded slice while keeping pressure zero. Lean proves:

1. the same strengthened pressure-seeded target still has complete regularity, dynamics, and energy certificates for this affine seed-blend candidate;
2. therefore the same-route version again reaches an explicit `AlmostFeffermanBridge` whose live mouth is compatibility;
3. a nearby descendant route with affine seed-blend compatibility closes immediately;

4. the same-route mouth already closes in special cases such as $(a, b) = (1, 0)$ and the zero-vorticity regime.

So the top-down NS tendril now reaches a second, strictly richer non-self frontier in Lean.

In `FeffermanSaturatedPressureFrontier.lean`, Lean adds the first nonlinear non-self candidate:

$$u(t, x) = a \cdot \frac{\omega(t, x)}{1 + |\omega(t, x)|}.$$

The file proves:

1. the saturated profile still satisfies the strengthened pressure-seeded target's regularity, dynamics, and energy slots;
2. the same-route version again reaches an explicit **AlmostFeffermanBridge** with compatibility as the only remaining mouth;
3. a nearby descendant route with saturated compatibility closes immediately;
4. zero vorticity already collapses the same-route mouth.

So the top-down ladder now contains not just linear and affine non-self frontiers, but also a first nonlinear saturated frontier in Lean.

In `FeffermanFrozenGatePressureFrontier.lean`, Lean adds a more manuscript-shaped nonlinear non-self candidate:

$$u(t, x) = a \cdot \frac{|\omega(1, x)|}{1 + |\omega(1, x)|} \cdot \omega(t, x).$$

The live tendril is now modulated by a nonlinear gate computed from the frozen seed slice. The file proves:

1. this frozen-seed-gated candidate still satisfies the strengthened pressure-seeded target's regularity, dynamics, and energy slots;
2. the same-route version again reaches an explicit **AlmostFeffermanBridge** with compatibility as the only remaining mouth;
3. a nearby descendant route with frozen-gate compatibility closes immediately;
4. zero vorticity again collapses the same-route mouth.

So the top-down ladder now contains not just linear, affine, and pointwise nonlinear frontiers, but also a first frozen-seed-gated frontier in Lean.

In `FeffermanSeedLiveProfileFrontier.lean`, Lean then packages these frontiers into one reusable theorem family. The new structure `SeedLiveProfile` covers any pointwise candidate of the form

$$u(t, x) = \Phi(\omega(1, x), \omega(t, x))$$

under the single bound

$$|\Phi(\text{seed}, \text{live})| \leq c(|\text{seed}| + |\text{live}|).$$

The file proves:

1. every such profile automatically carries the strengthened pressure-seeded target through regularity, dynamics, and energy;

2. the same-route version always reaches an explicit **AlmostFeffermanBridge** whose only remaining mouth is compatibility;
3. the descendant route with the exact profile-compatibility predicate closes immediately;
4. the scalar-rescaled, affine seed-blend, nonlinear saturated, and frozen-seed-gated candidates are recovered as concrete instances.

So the NS top-down lane now has not only several explicit non-self frontier examples, but also a reusable seed/live profile shell that subsumes them. Lean now also records the dual stationary closure regime already on that pointwise shell: if the profile returns the seeded slice pointwise and the concrete windowed vorticity tendril is stationary between the seeded slice and every live slice, then the same pressure-seeded clause closes directly there too.

In **FeffermanSeedLiveOperatorFrontier.lean**, Lean pushes this one step further to a genuinely non-pointwise shell. The new structure **SeedLiveOperator** allows the candidate at x to depend on the full seeded and live slices as functions, provided it satisfies a uniform envelope bound in terms of seeded/live slice bounds. The file proves:

1. every such operator automatically carries the strengthened pressure-seeded target through regularity, dynamics, and energy;
2. the same-route version again reaches an explicit **AlmostFeffermanBridge** whose only remaining mouth is compatibility;
3. the descendant route with exact operator compatibility closes immediately;
4. the pointwise **SeedLiveProfile** shell embeds into this operator shell, and a concrete anchored non-pointwise seed/live blend is realized as an instance.

So the top-down NS lane now reaches not only pointwise non-self frontier families, but also a first proof-carrying non-pointwise operator shell.

In **FeffermanSeedLiveSampleKernelFrontier.lean**, Lean specializes that operator shell to a finite sampled-kernel family. The new structure **SeedLiveSampleKernel** builds candidates of the form

$$u(t, x) = \sum_{i \in \kappa} \left(a_i \omega(1, \alpha_i(x)) + b_i \omega(t, \beta_i(x)) \right),$$

where the seed and live slices are sampled through finite anchor maps and the only global control is the total absolute kernel mass. The file proves:

1. every such finite sampled kernel automatically induces a **SeedLiveOperator**, so it inherits the same regularity, dynamics, and energy certificates for the strengthened pressure-seeded target;
2. the same-route version again reaches an explicit **AlmostFeffermanBridge** whose only remaining mouth is compatibility;
3. the descendant route with exact sampled-kernel compatibility closes immediately;
4. the new frontier is therefore a proof-carrying finite mixing shell sitting strictly between the generic pointwise profile family and the fully non-pointwise operator shell.

So the top-down NS lane now reaches not only general operator shells, but also a first finite sampled-kernel family that still lands at the same compatibility mouth.

In `FeffermanSeedLiveShiftKernelFrontier.lean`, Lean then imposes a first genuinely spatial structure on that sampled-kernel shell. The new structure `SeedLiveShiftKernel` specializes to translation-style kernels on an additive state space:

$$u(t, x) = \sum_{i \in \kappa} (a_i \omega(1, x + s_i) + b_i \omega(t, x + \ell_i)).$$

The file proves:

1. every such shift kernel canonically reduces to a `SeedLiveSampleKernel`, so it inherits the same regularity, dynamics, and energy certificates for the strengthened pressure-seeded target;
2. the same-route version again reaches an explicit `AlmostFeffermanBridge` whose only remaining mouth is compatibility;
3. the descendant route with exact shift-kernel compatibility closes immediately;
4. a concrete diagonal shift kernel is realized as an instance.

So the top-down NS lane now reaches not only arbitrary finite sampled mixtures, but also a first translation-kernel shell that begins to look like a genuine finite convolution front.

2.27 Concrete Windowed Descendant Route

The newest specialization step pushes the same top-down frontier back down to the actual concrete windowed heat package. In `WindowedCutoffApproximation.lean`, Lean proves that the genuinely windowed cutoff

$$u \mapsto u \cdot (\text{smoothTransition}(cu + 1) - \text{smoothTransition}(cu))$$

still passes through the current approximation shell. So leaving the stricter `EnergyProfile` shell is not yet a failure of truncation convergence or of the induced Cole–Hopf convergence.

In `WindowedColeHopfHeatPackage.lean`, Lean proves the quantitative window bound

$$|\chi_c(u)| \leq c^{-1},$$

which is enough to instantiate the present heat-decayed Cole–Hopf package. So the current concrete route already carries a genuine windowed shared approximation package and a genuine uniform vorticity tendril.

In `WindowedColeHopfHeatFoundationFrontier.lean` and `WindowedColeHopfHeatSeededModel.lean`, Lean then fixes the actual candidate fields on that concrete route:

1. velocity is the package vorticity field;
2. pressure is identically zero;
3. the candidate closes a stronger seeded local target with zero pressure, a distinguished initial slice at time 1, and a uniform pointwise bound.

So the current compatibility frontier is no longer only an abstract theorem schema. It now sits on one explicit concrete windowed route.

2.28 Concrete Windowed Same-Route Rigidity

The next new files show that the concrete route already has theorem-level same-route rigidity.

In `WindowedColeHopfHeatScaledSeededFrontier.lean`, Lean packages the scaled descendant $u(t, x) = a \omega(t, x)$ on the concrete windowed route. The new file `WindowedColeHopfHeatScaledPressureSeededFrontier.lean` then lifts that same concrete rescaled descendant to the stronger pressure-seeded target already used by the later non-self frontiers, producing explicit concrete `AlmostFeffermanBridge` and `TopDownFeffermanBridge` objects and the direct clause theorem `WeightedObservable.windowedColeHopfHeat_realizes_scaled_pressure_seeded_clause_of_zero_vorticity`. It now also records the two honest same-route closure regimes on that stronger concrete target: the rescaled descendant closes the original compatibility mouth when $a = 1$, and more generally it closes that mouth for any a when the underlying windowed vorticity tendril vanishes identically. Lean packages those special cases as direct clause theorems `WeightedObservable.windowedColeHopfHeat_realizes_scaled_pressure_seeded_clause_of_zero_vorticity` and `WeightedObservable.windowedColeHopfHeat_realizes_scaled_pressure_seeded_clause_of_zero_vorticity`. In `WindowedColeHopfHeatSelfCompatibilityObstruction.lean`, Lean proves that if this scaled descendant is forced back through the original self-compatibility mouth on a genuinely nonzero vorticity state, then the scaling must satisfy $a = 1$. Equivalently, if $a \neq 1$, then no bridge of that same-route type exists on a nonzero-vorticity state.

In `WindowedColeHopfHeatSeedBlendFrontier.lean`, Lean packages the affine seed-blend descendant

$$u(t, x) = a \omega(t, x) + b \omega(1, x).$$

The new file `WindowedColeHopfHeatFrozenGateFrontier.lean` adds the first seeded pointwise nonlinear modulation on the same concrete route:

$$u(t, x) = a \frac{|\omega(1, x)|}{1 + |\omega(1, x)|} \omega(t, x).$$

So the windowed route no longer reaches only scalar and affine descendants before entering the finite sampled-kernel ladder.

The follow-up file `WindowedColeHopfHeatFrozenGateObstruction.lean` already makes that frozen-gate mouth rigid. Lean proves that if original self-compatibility holds at a nonzero live witness, then

$$a \frac{|\omega(1, y)|}{1 + |\omega(1, y)|} = 1.$$

Since the frozen gate is always strictly less than 1, any such witness forces $a > 1$. So on any genuinely nonzero live state there can be no same-route bridge of this literal frozen-gate form when $a \leq 1$. Lean also proves that any two nonzero live witnesses on such a bridge must see the same frozen gate, so this first seeded nonlinear pointwise route is already rigid on the seed side too.

The new files `WindowedColeHopfHeatSeedLiveProfileFrontier.lean` and `WindowedColeHopfHeatSeedLiveProfileObstruction.lean` then package the shared pointwise mouth directly on the same concrete windowed route:

$$u(t, x) = P(\omega(1, x), \omega(t, x)).$$

The first file proves that the scaled, affine seed-blend, saturated, and frozen-gate descendants are exact specializations of this one generic shell. It also now exposes the honest same-route theorem surface for that concrete shell itself: if the profile already equals the live input pointwise, then the original compatibility mouth closes directly; and if the underlying windowed vorticity tendril vanishes identically and $P(0, 0) = 0$, then the same pressure-seeded clause closes again on that zero state. The same concrete file also exports the retained uniform-vorticity bound from the

descendant-route `TopDownFeffermanBridge`. The second proves the shared witness law: at any nonzero live witness,

$$P(\omega(1, y), \omega(t, y)) = \omega(t, y).$$

So the pointwise search space is now structurally organized. These descendants are no longer independent repair guesses; they all have to close through the same concrete pointwise compatibility equation on the actual windowed package.

The next new files, `WindowedColeHopfHeatSeedLiveOperatorFrontier.lean` and `WindowedColeHopfHeatSeedOnlyOperatorObstruction.lean` lift this one step further to the full non-pointwise operator shell:

$$u(t, x) = O(\omega(1, \cdot), \omega(t, \cdot), x).$$

The first file proves that the pointwise profile shell, the finite sampled-kernel shell, and the finite shift-kernel shell are all exact specializations of this same concrete operator frontier on the windowed route. It also now exposes the honest same-route theorem surface for the operator shell itself: if the operator already returns the live slice pointwise, then the original compatibility mouth closes directly; and if the underlying windowed vorticity tendril vanishes identically while the operator sends the zero seeded/live pair to zero, then the same pressure-seeded clause closes again on that zero state. The same concrete file also exports the retained uniform-vorticity bound from the descendant-route `TopDownFeffermanBridge`. These concrete operator theorems are backed by the corresponding generic route lemmas in `FeffermanSeedLiveOperatorFrontier.lean`. Lean now also records the dual stationary closure regime on that shell: if the operator returns the seeded slice pointwise and the concrete windowed vorticity tendril is stationary between the seeded slice and every live slice, then the same pressure-seeded clause closes again. So the non-pointwise operator frontier now has three honest same-route closure mouths at theorem surface: live-return, seed-return under stationarity, and zero-vorticity collapse. The second proves the shared same-route law:

$$O(\omega(1, \cdot), \omega(t, \cdot), y) = \omega(t, y).$$

So the non-pointwise search space is now structurally compressed as well. The sampled-kernel and shift-kernel descendants no longer sit as separate repair stories; they all have to close through the same concrete operator equation on the actual windowed package.

The next new file, `WindowedColeHopfHeatSeedOnlyOperatorObstruction.lean`, turns that shared operator mouth into the first genuinely dynamical obstruction. If the chosen operator ignores the live slice entirely, then same-route closure forces

$$\omega(t_1, y) = \omega(t_2, y)$$

for every witness y and every pair of times t_1, t_2 . So any witness where the concrete windowed route actually changes in time already rules out that whole seed-only non-pointwise family. This is the first explicit Lean result saying that a surviving repair must depend on the live state in a real way, not just on the frozen seeded slice. The same file immediately specializes this to finite sampled kernels with zero live coefficients, so the first finite non-pointwise shell already inherits this live-dependence requirement. Since finite shift kernels are only sampled kernels with translation anchors, the same obstruction now bites the first genuinely spatial seed-only shell too.

The next file, `WindowedColeHopfHeatLiveOnlySampleKernelObstruction.lean`, proves the dual live-dependent finite-kernel obstruction. If the sampled kernel ignores the seeded slice entirely, then same-route closure reduces to

$$\sum_i a_i \omega(t, \alpha_i(y)) = \omega(t, y).$$

So whenever all sampled live witnesses vanish but the target live value does not, the corresponding live-only sampled-kernel bridge cannot exist. Since shift kernels are just sampled kernels with translation anchors, the same obstruction now bites the first genuinely spatial live-only shell as well.

In `WindowedColeHopfHeatSampleKernelFrontier.lean`, Lean then proves that this concrete affine descendant is exactly the diagonal one-sample instance of the generic sampled-kernel shell. So the concrete windowed route now connects downward into the abstract ladder not only through the generic seed-blend frontier, but also through the finite sampled-kernel layer. This matters because the older affine same-route obstruction now survives inside a first finite non-pointwise family, not only inside the pointwise seed-blend shell. The same concrete file now also exposes the direct same-route theorem surface inherited from the operator shell: arbitrary sampled kernels admit a direct clause theorem once self-compatibility is supplied, the live-return regime closes that route immediately, the zero-vorticity regime closes that same-route mouth automatically, and the same sampled-kernel route now also records the seed-return stationary regime inherited from the operator shell. The descendant-route bridge still exports the original windowed uniform-vorticity envelope. In addition, the diagonal specialization now re-exports that finite-kernel route directly on the older seed-blend target: Lean packages a direct clause theorem for the diagonal sampled kernel, a seed-blend-self-compatible same-route closure theorem, a direct live-endpoint closure theorem at $(a,b)=(1,0)$, a direct stationary seed-return closure theorem at the $(a,b)=(0,1)$ seed endpoint, and the corresponding zero-vorticity closure theorem, all with the seed-blend initial slice as the visible boundary data.

In `WindowedColeHopfHeatShiftKernelFrontier.lean`, Lean pushes this one rung further upward into genuinely spatial descendants. On additive-monoid state spaces, the anchored blend

$$u(t, x) = a \omega(t, x) + b \omega(1, x + s) + d \omega(t, x + s)$$

is proved to be exactly a finite translation-kernel instance. So the concrete windowed route now reaches the abstract shift-kernel layer too, not merely the one-point sampled-kernel shell. The same concrete shift-kernel file now also packages the direct same-route clause theorem under self-compatibility, the inherited live-return regime, the automatic zero-vorticity closure on that shell, the inherited seed-return stationary regime, and the retained uniform-vorticity bound on the descendant-route bridge. The anchored translation specialization now also re-exports those same-route theorems on the anchored three-term descendant itself: Lean gives direct anchored-shift clause wrappers from anchored self-compatibility, from any live-return regime that makes the specialized operator collapse back to the live slice, from any seed-return stationary regime that makes the specialized operator collapse back to the seeded slice, and from the zero-vorticity regime.

The next file, `WindowedColeHopfHeatShiftKernelObstruction.lean`, makes that spatial mouth rigid on the same-route side. Lean proves three witness-forced coefficient laws for the anchored descendant:

1. if the shifted initial and shifted live slices vanish while the unshifted live slice is nonzero, then $a = 1$;
2. if the unshifted live and shifted live slices vanish while the shifted initial slice is nonzero, then $b = 0$;
3. if the unshifted live and shifted initial slices vanish while the shifted live slice is nonzero, then $d = 0$.

The matching bridge-level impossibility corollaries are formalized too. So the first concrete spatial descendant is not merely packaged in Lean; its original self-compatibility mouth is already sharply constrained there.

The next widening step, `WindowedColeHopfHeatMaskedShiftKernelObstruction.lean`, lifts that same rigidity from one anchored spatial descendant to a finite family. In that family, every seed term uses the same shift s , while each live term either stays at x or moves to $x + s$ according to a Boolean mask. Lean proves that the whole family collapses back to the three-coefficient anchored spatial route with coefficients given by the total unshifted-live mass, total seed mass, and total shifted-live mass. So the same witness patterns now force, at finite-family level,

$$a_{\text{base}} = 1, \quad b_{\text{seed}} = 0, \quad d_{\text{shift}} = 0,$$

with the corresponding bridge-level impossibility corollaries when any of these three identities is violated. So the first concrete spatial same-route rigidity is no longer a one-instance fact; it already controls a nontrivial finite spatial shell.

The next widening step, `WindowedColeHopfHeatGeneralShiftKernelObstruction.lean`, removes the common-shift restriction and works directly at the arbitrary finite translation-kernel shell. Lean proves three generic collapse laws there:

1. if every seeded anchor except 0 vanishes, every live anchor except 0 vanishes, and the unshifted initial witness is nonzero, then the total zero-anchor mass

$$a_{\text{seed},0} + a_{\text{live},0}$$

is forced to equal 1;

2. if every seeded anchor vanishes, every nonzero live anchor vanishes, and the unshifted live witness is nonzero, then the total zero-shift live mass is forced to equal 1;
3. if the live slice at the base point vanishes, every live anchor vanishes, every seeded anchor except a chosen shift u vanishes, and the u -shifted seeded witness is nonzero, then the total seeded mass at that anchor is forced to equal 0;
4. if the live slice at the base point vanishes, every seeded anchor vanishes, every live anchor except a chosen shift u vanishes, and the u -shifted live witness is nonzero, then the total live mass at that anchor is forced to equal 0.

So the linear same-route collapse is no longer confined to the anchored or masked common-shift subfamilies. It already lives at the full finite translation-kernel shell, stated in terms of anchorwise coefficient sums, with the matching bridge-level impossibility corollaries when any of those forced anchor masses is violated. In particular, the old identity-anchor sampled kernel law

$$a_{\text{seed}} + a_{\text{live}} = 1$$

is now visible as a full finite shift-kernel phenomenon at anchor 0, not as a separate low-level accident.

The follow-up file `WindowedColeHopfHeatGeneralShiftKernelMaskedSpecialization.lean` shows that the older masked common-shift shell is already recovered from this newer general linear theorem. In the nondegenerate case $s \neq 0$, Lean identifies the masked seed mass with the general seeded mass at anchor s , the masked unshifted-live mass with the general zero-shift live mass, and the masked shifted-live mass with the general live mass at anchor s . So the old masked coefficient laws and their impossibility corollaries are now derived from the full finite shift-kernel shell rather than reproved in parallel. The proof also keeps the exact caveat visible: when $s = 0$, the masked

shifted-live terms collapse onto the zero anchor, so the old unshifted and shifted live masses are no longer separated by the general shell.

The companion file `WindowedColeHopfHeatGeneralShiftKernelIdentitySampleSpecialization.lean` does the same job for the older identity-anchor sampled-kernel sum law. It packages any identity-anchor sampled kernel as the zero-shift specialization of the general linear shell, transports the top-down bridge across that exact initial-slice and candidate equality, and then derives the old law

$$a_{\text{seed}} + a_{\text{live}} = 1$$

from the full finite shift-kernel theorem itself. So even the identity-anchor sampled-kernel sum rule is no longer a parallel low-rung argument; it is now an explicit specialization of the top linear shell.

The new file `WindowedColeHopfHeatSampleKernelObstruction.lean` makes that survival explicit at bridge level on the diagonal one-sample kernel. The same witness patterns that forced $a + b = 1$, $a = 1$, or $b = 0$ for the pointwise affine descendant now force the same identities for a same-route bridge whose candidate is presented through the sampled-kernel shell. So the first finite non-pointwise family does not evade the affine same-route rigidity; it inherits it.

The next file `WindowedColeHopfHeatIdentitySampleKernelObstruction.lean` strengthens that conclusion from the one-point diagonal case to every finite sampled kernel whose seed and live anchors are both the identity. Lean proves that such a family collapses back to the affine seed-blend route with coefficients given by the total live and seed masses. So finite repetition at the same point still does not create a new same-route escape corridor; it merely repackages the old affine obstruction in a larger finite shell.

In `WindowedColeHopfHeatSeedBlendObstruction.lean`, Lean proves bridge-level rigidity laws for the same-route mouth:

1. a nonzero initial slice forces $a + b = 1$;
2. a witness with zero initial slice but nonzero live vorticity forces $a = 1$;
3. a witness with zero live vorticity but nonzero initial slice forces $b = 0$;
4. violating any of these identities rules out the corresponding same-route bridge on the concrete windowed route.

So the NS crux has now moved one step forward: the issue is no longer merely that abstract non-self frontiers exist. The issue is that on the concrete windowed route, original same-route closure is already provably rigid.

The newest nonlinear step adds `WindowedColeHopfHeatSaturatedFrontier.lean`, which packages the concrete saturated descendant

$$u(t, x) = a \frac{\omega(t, x)}{1 + |\omega(t, x)|}.$$

The corresponding obstruction file `WindowedColeHopfHeatSaturatedObstruction.lean` proves the first nonlinear same-route rigidity theorem on the concrete windowed route: if the saturated descendant is forced back through original self-compatibility at a nonzero vorticity witness, then

$$a = 1 + |\omega(t, x)|.$$

So in particular, on any nonzero-vorticity state there can be no same-route bridge of this saturated form when $a \leq 1$. The new bridge-level upgrade is stronger: any same-route saturated bridge now

forces all nonzero witnesses to have the same absolute vorticity magnitude. So two distinct nonzero magnitudes already rule out this nonlinear same-route closure entirely.

The next file `WindowedColeHopfHeatSaturatedSampleKernelObstruction.lean` proves that this nonlinear descendant already hits a boundary even at the diagonal one-sample sampled-kernel rung. If a sampled-kernel bridge of that form realizes the saturated candidate with $a \neq 0$, then all nonzero initial witnesses must have the same absolute vorticity magnitude. Therefore any state with two nonzero initial witnesses of distinct magnitudes rules out that same-route sampled-kernel realization altogether.

The new file `WindowedColeHopfHeatSaturatedIdentitySampleKernelObstruction.lean` pushes this one step beyond the diagonal case to every finite sampled kernel whose seed and live anchors are both the identity. Lean proves that if such a bridge realizes the saturated descendant with $a \neq 0$, then the same constant-magnitude rigidity still holds on the initial slice. So finite repetition at the same point does not open a nonlinear same-route escape corridor either.

The next file `WindowedColeHopfHeatSaturatedMaskedShiftKernelObstruction.lean` connects that nonlinear boundary to the widened finite spatial shell. Lean proves that if the saturated descendant is realized through the finite masked shift-kernel family and there is a witness with zero shifted initial slice, zero shifted live slice, and nonzero unshifted live slice, then the total unshifted-live mass satisfies

$$a_{\text{base}}(1 + |\omega(t, x)|) = a.$$

So if there are two such zero-shifted witnesses with distinct nonzero absolute vorticity magnitudes, the masked-shift shell cannot realize the saturated descendant at all. In other words, the nonlinear route already exits the widened finite spatial linear shell unless the surviving witnesses are constant-magnitude in exactly the same rigid way.

The next file `WindowedColeHopfHeatSaturatedShiftKernelObstruction.lean` pushes that boundary from the widened common-shift family to arbitrary finite translation kernels. Here the witness pattern is stated directly at the shell level: every seeded anchor must vanish, every nonzero live anchor must vanish, and the unshifted live witness must remain nonzero. Lean proves that under those hypotheses the total zero-shift live mass satisfies

$$a_0(1 + |\omega(t, x)|) = a.$$

The same file now also proves the nonlinear zero-anchor total-mass analogue of the newer linear top-shell theorem: if every nonzero seeded anchor and every nonzero live anchor vanish while the unshifted initial witness remains nonzero, then

$$(a_{\text{seed},0} + a_{\text{live},0})(1 + |\omega(1, x)|) = a.$$

So the old identity-anchor sampled-kernel saturated law is no longer just a lower-rung fact in waiting; the full nonlinear shift-kernel shell now already sees the exact zero-anchor total-mass quantity that an identity specialization would need. So if two such witnesses survive with distinct nonzero absolute magnitudes, the entire finite translation-kernel shell cannot realize the saturated descendant. This is the first NS point where the nonlinear boundary is no longer tied to a single concrete shell encoding; it already lives at the general finite shift-kernel level. The same Lean file now also specializes that obstruction back to the first anchored translation shell, so the original concrete spatial descendant does not reopen this nonlinear corridor either.

The companion file `WindowedColeHopfHeatSaturatedShiftKernelConstantMagnitudeFrontier.lean` records the first honest constructive survivor inside that same full nonlinear shell. If the concrete windowed vorticity tendril has constant absolute magnitude

$$|\omega(t, x)| = r$$

for every t and x , then Lean proves the saturated candidate is exactly the zero-shift live-only finite shift kernel with coefficient

$$\frac{a}{1+r}.$$

So in that degenerate constant-magnitude regime there is an explicit top-down bridge from the full finite shift-kernel shell back to the nonlinear saturated descendant itself. This is useful because it shows the nonlinear story is no longer purely negative: the current shell really does have a narrow surviving corridor, and the surrounding obstruction theorems are now cutting around that corridor rather than merely around empty space. The same file now also exports that survivor directly at theorem surface: under the same constant-magnitude hypothesis, Lean packages the resulting saturated clause as a direct corollary rather than leaving the route only in existential bridge form.

The new file `WindowedColeHopfHeatSaturatedShiftKernelInvariantFrontier.lean` writes down the first honest widening of that positive corridor. Lean now proves that if the windowed vorticity tendril is stationary between the seeded slice and the live slice, invariant under every seed and live shift used by a finite kernel, and still has constant absolute magnitude r , then any finite shift kernel with total coefficient mass

$$a_{\text{seed}} + a_{\text{live}} = \frac{a}{1+r}$$

realizes the saturated candidate. So the nonlinear full shift-kernel shell is not merely nonempty at one hand-picked zero-shift point: it contains a whole invariant finite-kernel corridor. That is still highly rigid, but it is a real constructive shell theorem rather than just a surviving singleton. Lean now also exposes the same corridor directly on the nonlinear saturated target: once the finite kernel has the right total mass and the concrete windowed tendril is shift-invariant with constant magnitude, the saturated clause follows immediately via that finite shift-kernel route.

The companion file `WindowedColeHopfHeatSaturatedShiftKernelInvariantObstruction.lean` now closes the same shell at the candidate-equality level. Under the same stationarity and the same invariance under every seed and live shift used by the kernel, if the finite shift-kernel candidate coincides with the nonlinear saturated descendant and the total coefficient mass is nonzero, then every two nonzero witnesses must have the same absolute magnitude. So the full invariant finite-kernel corridor is no longer only constructive; it is already exact in the same constant-magnitude sense.

That file now also contains the corresponding proof-carrying bridge-level corollary. If a top-down bridge through the same finite invariant shift-kernel shell has candidate equal to the nonlinear saturated descendant, then the same equal-absolute-magnitude conclusion follows; equivalently, unequal nonzero witness magnitudes rule out the existence of such a bridge. So the full invariant shell now has both sides in Lean: a constructive corridor and a matching $\neg\exists B$ obstruction.

The next file `WindowedColeHopfHeatSaturatedAnchoredShiftInvariantFrontier.lean` pushes that positive corridor back down to the first concrete spatial anchored descendant

$$b\omega(t, x) + c\omega(1, x + s) + d\omega(t, x + s).$$

Under the same stationarity, the same invariance under $x \mapsto x + s$, the same constant-magnitude regime $|\omega(t, x)| = r$, and the concrete mass law

$$b + c + d = \frac{a}{1+r},$$

Lean proves that this anchored spatial candidate is exactly the nonlinear saturated descendant itself, and therefore yields a concrete top-down bridge through the anchored translation shell. So the new positive corridor is no longer merely abstract full-shell data: it already reaches a genuinely

spatial windowed descendant. The same file now also exports the direct clause-level corollary on the saturated target, so the anchored spatial corridor is usable without unpacking an existential bridge witness first.

The companion file `WindowedColeHopfHeatSaturatedAnchoredShiftInvariantObstruction.lean` then proves that this concrete anchored corridor is already exact at the candidate-equality level once the total coefficient sum $b + c + d$ is nonzero. Under the same stationarity and the same invariance under $x \mapsto x + s$, if the anchored spatial candidate coincides with the nonlinear saturated descendant, then every two nonzero witnesses must have the same absolute magnitude. So unequal nonzero magnitudes do not just block the older sampled or masked shells; they already rule out exact equality for this first concrete spatial survivor itself.

The follow-up file `WindowedColeHopfHeatSaturatedShiftKernelInvariantAnchoredSpecialization.lean` then closes the hierarchy gap: it shows that this anchored exactness theorem is not a parallel argument, but an explicit specialization of the newer full invariant finite shift-kernel exactness theorem through the anchored translation-kernel transport. So at the nonlinear exactness level the ladder is now cleaner too: full invariant shell first, anchored spatial theorem second.

The follow-up file `WindowedColeHopfHeatSaturatedShiftKernelMaskedSpecialization.lean` then shows that the earlier masked common-shift shell really is a special case of this newer general theorem. In the nondegenerate case $s \neq 0$, Lean identifies the old masked unshifted-live coefficient with the new shell-level zero-shift live mass, and then derives the old masked coefficient law and its impossibility corollary directly from the general finite shift-kernel obstruction. The proof also exposes the exact caveat: when $s = 0$, the “shifted” live terms collapse back onto the zero anchor, so the old masked base mass and the new shell-level zero-shift mass are no longer the same quantity.

The positive side now exists there too. The new file `WindowedColeHopfHeatSaturatedMaskedShiftKernelFrom` packages the constructive corridor on that masked route directly at theorem surface: if the total coefficient mass of the masked common-shift kernel is

$$\frac{a}{1 + r},$$

the concrete windowed vorticity tendril is invariant under the single common shift $x \mapsto x + s$, and $|\omega(t, x)| = r$, then the masked route already realizes the nonlinear saturated clause. With the additional stationarity hypothesis between the seeded and live slices, Lean upgrades the same route to a proof-carrying top-down bridge whose candidate is exactly the nonlinear saturated descendant. So the masked specialization is no longer only an obstruction corollary of the full shell: its constant-magnitude survivor is now exposed directly too.

The next file `WindowedColeHopfHeatSaturatedShiftKernelIdentitySampleSpecialization.lean` does the nonlinear identity-anchor job as well. It packages any identity-anchor sampled kernel as the zero-shift specialization of the full nonlinear shift-kernel shell, transports the top-down bridge across the exact initial-slice equality, and then derives the old identity-anchor saturated law

$$(a_{\text{seed}} + a_{\text{live}})(1 + |\omega(1, x)|) = a$$

and the corresponding constant-magnitude impossibility corollary directly from the top nonlinear shell. So even the older identity-anchor sampled-kernel saturated rigidity is no longer a parallel low-rung argument; it is now an explicit specialization of the full finite nonlinear shift-kernel theorem.

That route now also has the matching constructive theorem surface. The file `WindowedColeHopfHeatSaturatedI` proves that if both anchor families are the identity, the total seed/live mass is

$$a_{\text{seed}} + a_{\text{live}} = \frac{a}{1 + r},$$

and the concrete windowed vorticity tendril has constant absolute magnitude $|\omega(t, x)| = r$, then the finite sampled-kernel route already realizes the nonlinear saturated initial slice itself. With the additional stationarity hypothesis between the seeded slice and the live slice, Lean upgrades that specialization to a proof-carrying sampled-kernel bridge whose candidate is exactly the nonlinear saturated descendant. So the identity-anchor shell now has both sides too: the old impossibility law remains, but the constant- magnitude corridor is now exposed directly at theorem surface rather than only through the ambient shift-kernel file.

The diagonal one-sample shell now has that same positive nonlinear surface too. The new file `WindowedColeHopfHeatSaturatedDiagonalSampleKernelFrontier.lean` simply instantiates the identity-anchor corridor on `diagonalSampleKernel`: if

$$p + q = \frac{a}{1 + r}$$

and the concrete windowed vorticity tendril has constant absolute magnitude $|\omega(t, x)| = r$, then the diagonal sampled route already realizes the nonlinear saturated clause; with stationarity, Lean upgrades it to a proof-carrying diagonal sampled-kernel bridge whose candidate is exactly the nonlinear saturated descendant. So the old affine one-point shell now has a direct constructive saturated wrapper in addition to its exactness laws.

3 Bounded-Energy Finite-Mode Schwartz Cell

The finite-mode linear and matrix-linear whole-space candidates were useful as obstruction tests, but their spatial profiles grow at infinity and Lean now formally rejects them from the finite-energy witness class. The positive replacement is to keep the finite-dimensional mode idea while changing the spatial basis: `NavierStokesFiniteModeBoundedEnergy.lean` introduces `twoModeSchwartzVelocity`

$$u(t, x) = a(t)f(x) + b(t)g(x),$$

where $f, g \in \mathcal{S}(\mathbb{R}^3, \mathbb{R}^3)$, together with the initial slice `twoModeSchwartzInitialVelocity` and the pressure class `affineAddScalarSchwartzPressure`. Lean proves that this two-mode Schwartz ansatz is finite-energy in the precise local sense whenever the scalar amplitudes are uniformly bounded: `finiteInitialKineticEnergy_twoModeSchwartzInitialVelocity`, `MatchesInitialVelocity_twoModeSchwartz`, `boundedKineticEnergy_twoModeSchwartzVelocity_of_abs_le`, and `boundedKineticEnergyUpTo_twoModeSchwartz`. The packaged theorem `twoModeSchwartzVelocity_finiteInitial_matches_boundedEnergyUpTo_of_abs_le` records the exact constructive status: bounded-energy finite-mode exploration is now possible without leaving the whole-space finite-energy domain.

The latest obstruction pass makes the old matrix-linear boundary exact at the initial-energy interface. In `NavierStokesFiniteModeEnergyObstruction.lean`, Lean proves `finiteInitialKineticEnergy_finite` the time-zero slice $x \mapsto g(0)Ax$ has finite initial kinetic energy if and only if either every entry of A is zero or $g(0) = 0$. Thus the whole-space matrix-linear ansatz cannot be repaired into the Clay-style finite-energy input class while it remains an active nonzero linear-growth profile; a repair must genuinely localize or decay the spatial modes.

The same boundary now holds at the finite-time witness layer on nonnegative slabs. Lean proves `nonempty_ExplicitFiniteTimeRegularityWitness_finiteModeLinearMatrixTimeVelocity_initial_iff`: when $0 \leq T$, the witness type for initial datum $x \mapsto g(0)Ax$ is inhabited if and only if $A = 0$ entrywise or $g(0) = 0$. The forward direction uses the general theorem that every witness on a nonnegative slab has finite initial energy; the reverse direction is the canonical zero witness after the degenerate initial slice is identified with zero.

The slab-bounded-energy boundary is now exact as well. Lean proves `boundedKineticEnergyUpTo_finiteModeL` on a fixed slab $0 \leq t \leq T$, the velocity $u(t, x) = g(t)Ax$ satisfies the finite-time bounded-energy predicate if and only if either $A = 0$ entrywise or $g(t) = 0$ throughout that slab. Thus the obstruction is not an artifact of the time-zero input filter; any nonzero matrix-linear mode that is active even once on the slab is excluded directly by the finite-time energy slot.

The finite-time smooth PDE package now has the same exact boundary. Lean proves `finiteModeLinearMatrix_t` for symmetric trace-free A , smooth g, g' , and a derivative relation, the affine-quadratic pressure package gives smooth (u, p) , pointwise momentum, and incompressibility on the slab $0 \leq t \leq T$, while `boundedKineticEnergyUpTo` for u is equivalent to $A = 0$ entrywise or $g(t) = 0$ on the whole slab. Thus the slab theorem is not just a predicate-level obstruction; it classifies the energy slot inside the concrete finite-time PDE package.

The global bounded-energy boundary is exact too. Lean proves `boundedKineticEnergy_finiteModeLinearMatr` the whole-time predicate `boundedKineticEnergy` holds for $u(t, x) = g(t)Ax$ if and only if either $A = 0$ entrywise or $g(t) = 0$ for every time $t \in \mathbb{R}$. This rules out repairs that preserve a nonzero linear-growth matrix mode but try to move its activity outside the chosen slab; the global finite-energy slot sees every active time slice.

The exact boundary now reaches the whole smooth PDE package rather than only the velocity predicate. Lean proves `finiteModeLinearMatrix_timeDependent_global_smooth_momentum_incompressible_b` for symmetric trace-free A , smooth g, g' , and the derivative relation between them, the affine-quadratic pressure construction yields smooth (u, p) , pointwise momentum, and incompressibility at all times, while `boundedKineticEnergy` for the constructed u is equivalent to $A = 0$ entrywise or $g \equiv 0$. This is still not analytic existence progress; it is an exact incompatibility certificate for the polynomial matrix-linear finite-mode route.

The same file now also discharges the basic smoothness side conditions for this cell. The theorem `smoothSpaceTimeVelocity_twoModeSchwartzVelocity` proves that smooth scalar amplitudes a, b and Schwartz spatial profiles f, g give a smooth velocity field on $\mathbb{R} \times \mathbb{R}^3$, while `smoothSpaceTimePressure_affineAddScalarSchwartzPressure` packages the corresponding affine-plus-localized Schwartz pressure smoothness. The stronger package `twoModeSchwartzVelocity_finiteInitial_sm` therefore gives finite initial energy, smoothness, initial matching, and slab-bounded energy from smooth bounded amplitudes.

The next Lean pass also closes the incompressibility side condition for this finite-mode cell. The theorem `spatialDivergence_twoModeSchwartzVelocity` identifies the spatial divergence of $a(t)f(x) + b(t)g(x)$ as the amplitude-weighted sum of the initial divergences of f and g . Consequently `spatialDivergence_twoModeSchwartzVelocity_zero` proves that divergence-free Schwartz profiles produce an incompressible two-mode space-time velocity. The package `twoModeSchwartzVelocity_finiteInit` now records finite initial energy, smoothness, initial matching, incompressibility, and slab-bounded energy together.

This is not yet a Navier–Stokes existence theorem for arbitrary Schwartz data. The constructor `ExplicitFiniteTimeRegularityWitness.of_twoModeSchwartzVelocity_affineAddScalarSchwartzPressure` only produces a finite-time witness after the smoothness, momentum, incompressibility, and initial-condition hypotheses are supplied. Its value is that the energy slot is no longer the obstruction for this finite-mode family; the remaining burden is the real PDE closure of the chosen modes and pressure. The companion constructor `ExplicitFiniteTimeRegularityWitness.of_twoModeSchwartzVelocity_affineAdd` removes the smoothness hypotheses from that witness interface by deriving them from coefficient smoothness, leaving the pointwise momentum equation, incompressibility, initial match, viscosity sign, derivative witnesses, and amplitude bounds as the explicit remaining assumptions. The stronger constructor `ExplicitFiniteTimeRegularityWitness.of_twoModeSchwartzVelocity_affineAddScalarSchwartzPr` also derives incompressibility from the two profile-level divergence-free assumptions. At that point

the only genuinely PDE-specific remaining assumption in this finite-mode witness lane is the pointwise momentum equation, together with bookkeeping data for the initial match, viscosity, scalar derivatives, and amplitude bounds.

The next time-dependent cleanup removes two of those bookkeeping obligations from the typed two-mode surface. Lean now proves `hasDerivAt_deriv_of_contDiff`, extracting `HasDerivAt a (deriv a(t)) t` from a smooth scalar amplitude, and `timeVelocityDerivative_twoModeSchwartzVelocity_deriv`, identifying the concrete time derivative of $a(t)f(x) + b(t)g(x)$ as $(\text{deriv } a(t))f(x) + (\text{deriv } b(t))g(x)$. The constructor `ExplicitFiniteTimeRegularityWitness.of_twoModeSchwartzVelocity_affineAddScalarSch` then uses those derivative witnesses and the canonical initial slice internally. Thus, on the smooth divergence-free bounded two-mode Schwartz surface, the external finite-time witness inputs have been narrowed to the pointwise momentum equation, viscosity sign, and amplitude bounds, plus the smoothness and profile divergence-free assumptions that define the cell.

The latest Lean pass closes that momentum slot for a concrete finite-energy member of the same Schwartz family. For the zero-amplitude specialization $a = b = 0$, Lean proves `twoModeSchwartzVelocity_zero_zero`, so the velocity is literally the zero velocity field independently of the Schwartz profiles f, g . It also proves `affineAddScalarSchwartzPressure_zero_timeOnly`, identifying the pressure class with a time-only gauge when the affine spatial coefficient and localized Schwartz amplitude vanish. Using the existing zero-velocity Navier–Stokes theorem, this yields `momentumEquation_zeroTwoModeSchwartzVelocity_affineTimeOnlyPressure`, an actual pointwise proof of

$$\partial_t u + (u \cdot \nabla)u + \nabla p = \nu \Delta u$$

inside the bounded-energy two-mode Schwartz construction. This is still the degenerate zero-amplitude case rather than a nonzero finite-dimensional Schwartz heat closure, but it removes the previous external `heq` placeholder for at least one genuine Schwartz member of the cell.

The follow-up Lean pass moves the momentum work to genuinely nonzero two-mode amplitudes. Lean now proves `spatialLaplacian_twoModeSchwartzVelocity`, expanding Δu modewise, and `spatialConvection_twoModeSchwartzVelocity`, expanding $(u \cdot \nabla)u$ into the two self-interaction terms and the two mixed directional-derivative terms. Combining those expansions with the time derivative identity yields `momentumEquation_twoModeSchwartzVelocity_of_explicitClosure`: for arbitrary smooth amplitudes a, b , the pointwise Navier–Stokes momentum equation follows once the affine-plus-Schwartz pressure gradient closes the fully expanded finite-mode residual. The specialization `momentumEquation_oneOneTwoModeSchwartzVelocity_of_explicitClosure` records the constant nonzero-amplitude case $a(t) = b(t) = 1$, and `oneOneTwoModeSchwartzVelocity_nonzero_of_exists_pro` shows that this velocity is genuinely nonzero whenever the two profiles do not cancel everywhere. This still leaves the harder constructive task of producing a nonzero Schwartz profile pair and affine-plus-Schwartz pressure satisfying the explicit residual closure, but the PDE momentum proof is no longer the zero-field tautology and no longer hides the nonlinear term.

The next Lean pass specializes this nonzero-amplitude momentum result to the inviscid zero-pressure boundary case. The theorem `momentumEquation_oneOneTwoModeSchwartzVelocity_inviscid_zeroPres` states that, for $a(t) = b(t) = 1$ and $\nu = 0$, the pointwise equation

$$\partial_t u + (u \cdot \nabla)u + \nabla p = 0 \Delta u$$

with $p = 0$ follows directly from the vanishing of the explicit self-plus-mixed convection expansion. Lean also packages this as `ExplicitFiniteTimeRegularityWitness.of_oneOneTwoModeSchwartzVelocity_invi`, which discharges smoothness, initial matching, incompressibility, bounded energy, amplitude bounds, and viscosity sign for the constant-one two-mode Schwartz cell. Finally, `oneOneTwoModeSchwartzVelocity_nonzero`

records the nonzero side condition via $\exists x, f(x)+g(x) \neq 0$. This is deliberately weaker than a positive-viscosity Navier–Stokes construction: it isolates a usable exact nonzero finite-energy momentum witness at the Euler boundary while keeping the profile-level convection closure explicit.

The current Lean pass then returns to positive viscosity and constructs a concrete residual-closure baseline with nonzero Schwartz profiles. It imports the Mathlib smooth bump-function API and defines `nsUnitBumpSchwartzVelocity`, a compactly supported smooth first-coordinate bump converted to a Schwartz velocity profile. Lean proves `nsUnitBumpSchwartzVelocity_nonzero` by evaluating the bump at the origin. Pairing this profile with its negative gives the anti-profile identity `oneOneAntiProfileSchwartzVelocity_zero`; hence the total two-mode velocity cancels even though both spatial profiles are nonzero. Using the zero-velocity momentum theorem, Lean proves `momentumEquation_oneOneAntiProfileSchwartzVelocity_affineTimeOnlyPressure_of_posViscosity` and then expands it back to the profile-level residual in `explicitClosure_oneOneAntiProfileSchwartzVelocity_`. The resulting existential theorem `exists_nonzeroSchwartzProfiles_pressure_explicitClosure_fullViscous` states that, for every $\nu > 0$, there exist nonzero Schwartz profiles f, g and a pressure field p (in the proof, $p = 0$) satisfying the fully expanded viscous residual closure. This is an exact Lean construction, but it is still a cancellation construction: it does not yet provide a non-cancelling positive-viscosity finite-energy flow.

4 Current Boundary

The following are *not* yet formalized in Lean in the present NS lane:

1. the manuscript’s specific chart-cutoff datum on SDiff^m ;
2. the finite-mode truncation / approximation theorem actually needed for that datum;
3. instantiations of the new top-down bridge layers with real velocity and pressure candidates, real regularity and dynamics certificates, and a real vorticity-to-velocity compatibility theorem beyond the current self, scalar-rescaled, affine seed-blend, nonlinear saturated, frozen-seed-gated, bounded seed/live operator toy candidates, and the current concrete windowed descendants;
4. strengthening the current seeded / bounded-seeded self models toward something recognizably closer to the Clay/Fefferman predicates than pointwise bounds, zero pressure, and a distinguished initial slice;
5. the local hypotheses $(H1)$, $(H2)$;
6. the manuscript’s claimed identity-only propagation lemma;
7. a non-cancelling positive-viscosity two-mode Schwartz profile and pressure satisfying the expanded residual closure;
8. the BKM-closing bridge for the full infinite-dimensional argument.

So the honest current status is:

Lean has formalized the generic interface shell that a proof attempt would need, but not the manuscript-specific analytic instantiations.

5 Verification

The newest invariant-shell exactness step was checked directly with:

```
lake build Mettapedia.FluidDynamics.NavierStokes.\
WindowedColeHopfHeatSaturatedShiftKernelInvariantObstruction
lake build Mettapedia.FluidDynamics.NavierStokes.\
WindowedColeHopfHeatSaturatedAnchoredShiftInvariantFrontier
lake build Mettapedia.FluidDynamics.NavierStokes.\
WindowedColeHopfHeatSaturatedAnchoredShiftInvariantObstruction
lake build Mettapedia.FluidDynamics.NavierStokes.\
WindowedColeHopfHeatSaturatedShiftKernelInvariantAnchoredSpecialization
```

The newest bounded-energy finite-mode Schwartz step was checked directly with:

```
lake build Mettapedia.FluidDynamics.NavierStokes.\
NavierStokesFiniteModeBoundedEnergy
lake build Mettapedia.FluidDynamics.NavierStokes.\
NavierStokesFiniteModeBoundedEnergyRegression
```

The full active NavierStokes lane was then rechecked with the grouped build command

```
mods=$(find Mettapedia/FluidDynamics/NavierStokes -maxdepth 1 -name '*.lean' \
  | sort | sed 's|/|.lg; s|\.lean$||')
lake build $mods
```

6 Next Lean Target

The natural next formal step is no longer more prose. It is to keep pushing top-down from the target:

1. strengthen the proof-carrying target again by making at least one more regularity or dynamics predicate genuinely nontrivial while keeping the clause closable by current packages;
2. push the concrete windowed route beyond same-route rigidity and toward a real velocity-generation or pressure-generation theorem on that fixed candidate;
3. reconnect whichever concrete descendant survives that test to the mode-space chart / approximation shells.